

CP References — ICPC Team Notebook

Contest Notes	2
Setup & don't-forget	2
Number theory	2
Bits	2
Games (Grundy)	2
Expected value	2
Misc	2
Tricks	3
Number Theory	4
Divisors	5
Euler's Phi (Totient)	6
Extended GCD & CRT	7
Linear Sieve	8
Miller-Rabin Primality	9
Sieve of Eratosthenes	10
Combinatorics	11
Combinatorics Toolkit	12
Partitions (Pentagonal Theorem)	13
nCr - Pascal's Triangle	14
DP	15
Digit DP	16
Longest Increasing Subsequence (LIS)	17
Matrix Exponentiation	18
Rerooting	19
SOS DP (Sum over Subsets)	20
Tree Knapsack	21
Graphs & Trees	22
Articulation Points	23
Bellman-Ford	24
Bridges (Tarjan)	25
Bridge Tree	26
Dijkstra	27
DSU (Union-Find)	28
DSU with Rollback	29
Floyd-Warshall	30
Heavy-Light Decomposition (HLD)	31
LCA (Binary Lifting)	33
MST (Kruskal)	34
Strongly Connected Components (SCC)	35
Topological Sort	36
Range Queries	37
Fenwick Tree (Point Update, Range Query)	38
Fenwick Tree (Range Update, Point Query)	39
Lazy Segment Tree	40
Mo's Algorithm	42
Mo on Tree	43
Mo with Rollback	45
Segment Tree	47
Segment Tree (Iterative)	48
Sliding Window Min/Max	49
Sparse Table	50
Sqrt Decomposition	51
Strings	52
Binary Trie (XOR)	53
String Hashing	54
KMP (Prefix Function)	55
Manacher (Palindromes)	56
Multiset Hashing	57
Trie	58
Geometry	59
Closest Pair of Points	60
Geometry Basics	61
Other	62
Bit Tricks	63
Coordinate Compression	64
Ordered Multiset (PBDS)	65
Random (RNG)	66

Contest Notes

Quick facts that solve or unlock problems. (Code lives in the folders; see README index.)

Setup & don't-forget

- **Fast IO** — first lines of main: `ios_base::sync_with_stdio(false); cin.tie(nullptr);` — after this NEVER mix in `scanf/printf`.
- `'\n'` not `endl` (`endl` flushes = slow in loops). **Exception: interactive problems** — there you MUST flush every query (`endl` or `cout.flush()`).
- **Interactive protocol**: `cout << "? " << x << endl`; then read the reply — `endl` everywhere, no buffering games. Note `cin.tie(nullptr)` disables the auto-flush-before-read, so with fast IO every un-flushed query = guaranteed idleness TLE.
- If the judge replies `-1` (or anything out of protocol) $\rightarrow$ `exit(0)` **immediately**: you exceeded the query budget or sent garbage, and looping on turns a clean WA into a confusing TLE.
- Finish with `! answer + flush`; count queries against the budget first (it's usually $\sim \log_2 n$ — an off-by-one loop bound is the classic interactive WA).
- Multi-test skeleton: `int t; cin >> t; while (t--) solve();` — **reset globals/arrays inside solve()**, half of all WAs on test 2 are stale state.
- Real-number answers: `cout << fixed << setprecision(10)`; — default prints 6 significant digits and truncates big values into scientific notation.
- Overflow: `1e5 * 1e5` overflows `int` **silently** — cast first: `(1ll)a * b`. Shifts: `1LL << k`, not `1 << k`, for $k \geq 31$. Constants: `const ll INF = 4e18;` (`int` max $\sim 2.1e9$).
- Big arrays **global**, never inside main — local ones overflow the stack. Same for recursion deeper than $\sim 1e5$ on some judges.
- `v.size()` is unsigned: `v.size() - 1` on an empty vector = huge number — write `(int)v.size()`.
- `sort` comparator must be **strict** (`<`, never `<=`) — a non-strict one is undefined behavior (can RE, not just mis-sort).
- Read until EOF: `while (cin >> x) { ... }`. File IO when the judge wants it: `freopen("in.txt", "r", stdin); freopen("out.txt", "w", stdout);`
- Compile: `g++ -O2 -std=c++17 a.cpp -o a` — while debugging add `-g -fsanitize=address,undefined` (catches overflow, out-of-bounds, UB at runtime).

Number theory

- The number of divisors of N is **odd iff N is a perfect square** (divisors pair up except $\sqrt{N}$). $\rightarrow$ `Number Theory/divisors.cpp`
- All common divisors of a, b are exactly the divisors of **`gcd(a, b)`**.
- Starting from 0 and repeatedly adding `STEP mod MOD`, you can reach **every** value $0..MOD-1$ **iff `gcd(STEP, MOD) = 1`**.
- Every even number > 2 is a sum of two primes (Goldbach — a conjecture, but verified far past any contest limit).
- The largest gap between consecutive primes below $1e9$ is only **~ 300** $\rightarrow$ "find a prime near x " loops terminate almost instantly.
- N is divisible by 6 iff N is even **and** its digit sum is divisible by 3.
- Euler's theorem: **$a^{\phi(m)} \equiv 1 \pmod{m}$** when $\gcd(a, m) = 1$.
- Huge exponents: **$a^n \equiv a^{(n \bmod \phi(m) + \phi(m))} \pmod{m}$** for $n \geq \log_2(m)$ — safe even when $\gcd(a, m) \neq 1$; drop the $+\phi(m)$ only if $\gcd = 1$. $\rightarrow$ `Number Theory/euler-phi.cpp`

Bits

- `n & (n-1)` — turn off the rightmost set bit (also: loop it to count set bits / test power of two).
- `n & (n+1)` — clear all trailing ones.
- `n | (n+1)` — set the rightmost zero bit.
- `n & -n` — extract the lowest set bit; note it always **divides n** .
- Number of set bits in $[1, 2^k - 1] = k \cdot 2^{k-1}$.
- Count of numbers in $[l, r]$ with a given bit set: $O(1)$ per bit. $\rightarrow$ `Other/bit-tricks.cpp`
- Prefix XOR $1^2^3 \dots^n$ is periodic mod 4: **$n, 1, n+1, 0$** for $n \% 4 = 0, 1, 2, 3$. $\rightarrow$ `Other/bit-tricks.cpp`
- N is a sum of K powers of two (repeats allowed) **iff `popcount(N)  $\leq K \leq N$`** .

Games (Grundy)

- Grundy number of a position = **mex of the Grundy numbers of positions reachable in one move**. Position is LOSING for the player to move **iff $g = 0$** . Compute by memoized DFS over the game graph.
- **Sum of independent games** (each turn: move in exactly ONE component): total $g =$ **XOR of the components' g** — losing iff the XOR is 0. This is why single-game Grundy tables solve multi-pile problems.
- **Nim** (take any amount from one pile): $g(\text{pile of } n) = n \rightarrow$ first player wins iff XOR of piles $\neq 0$. Winning move: pick a pile where $p \wedge \text{total} < p$, shrink it to $p \wedge \text{total}$.
- Subtraction game (take $1..k$): **$g(n) = n \bmod (k+1)$** . Unknown small game? **Brute-force $g(n)$ for $n \leq 60$ and look for the period** — most simple games are eventually periodic.
- **Misère nim** (who takes last LOSES): play exactly like normal nim, EXCEPT when every pile has size 1 — then first wins iff the number of piles is **even**.
- **Staircase nim** (move stones one step down a staircase): only stones on **odd steps** matter — XOR those pile sizes.
- No Grundy needed if you can find a **mirroring/pairing strategy**: copy the opponent's move in the symmetric half — proves a win without any computation. Also try "strategy stealing" for existence proofs.

Expected value

- **Linearity of expectation**: $E[X+Y] = E[X] + E[Y]$ — **works even when dependent**. The main move: write the target as a sum of 0/1 indicators, then $E[\text{total}] = \sum P(\text{indicator} = 1)$. ("Expected # of visible towers / inversions / adjacent equal pairs...")
- Nonnegative integer X : **$E[X] = \sum P(X \geq k)$, $k = 1, 2, \dots$** — often much easier than finding the full distribution ("expected number of rounds survived").
- Waiting times: success probability $p \rightarrow$ expected tries = **$1/p$** . **Coupon collector**: expected draws to see all n types = $n \cdot (1 + 1/2 + \dots + 1/n) \approx n \ln n$.
- k independent uniform $[0, 1]$: **$E[\min] = 1/(k+1)$, $E[\max] = k/(k+1)$** , i -th smallest = $i/(k+1)$. Same fractions for " k random cuts of a stick".
- Random permutation of n : expected **fixed points** = **1** (any $n!$), expected **# cycles** = **$1 + 1/2 + \dots + 1/n$** , $P(\text{two given elements share a cycle}) = 1/2$.
- Random process with states $\rightarrow$ set unknowns $E[\text{state}] =$ expected steps to finish, write one linear equation per state, solve: telescoping for chains/lines, **Gaussian elimination** for $\leq$ a few hundred states.
- **Gambler's ruin** (fair ± 1 walk from i , absorbing at 0 and n): $P(\text{reach } n \text{ first}) = i/n$; expected steps = **$i \cdot (n-i)$** .

Misc

- Chessboard $N \times M$: number of white squares = **$(N \cdot M + fg) / 2$** where $fg = 1$ if the corner square is white, else 0.

- A binary string with no palindromic substring of length > 1 has length ≤ 2 (adjacent chars must differ AND $s[i] \neq s[i+2]$ — contradiction from length 3 on).
- Any swap in a permutation changes the inversion count by an **odd** amount (flips parity). So 3-cycles preserve parity, and a duplicate value works as a "free parity reset" (swapping equal elements changes nothing).
- Minimum swaps to sort a permutation = $n - (\text{number of cycles in it})$.

Tricks

"See X $\rightarrow$ reach for Y" recipes. Skim when a problem feels unapproachable.

Picking the attack

- "**Minimize the maximum**" / "maximize the minimum" / "k-th smallest value" $\rightarrow$ **binary search on the answer**: write `possible(x)`, it's monotonic; for k-th smallest, count how many $\leq \text{mid}$.
- "Count subarrays with $\text{sum/xor} = k$ " $\rightarrow$ **prefix + hashmap**: $\text{pre}[r] - \text{pre}[l-1] = k$ means look up $\text{pre}[r] - k$ among earlier prefixes (xor: $\text{pre}[r] \oplus k$).
- Window property **monotonic** (adding elements only hurts / only helps) $\rightarrow$ **two pointers** instead of binary search per position.
- "Sum of f over ALL subarrays/pairs" $\rightarrow$ flip to **per-element contribution**: $a[i]$ sits in $(i+1) \cdot (n-i)$ subarrays; count how often each element/bit/pair is counted instead of iterating objects.
- Same flip for **subsets**: each element lies in 2^{n-1} of the 2^n subsets $\rightarrow$ sum of sizes over all subsets = $n \cdot 2^{n-1}$; each PAIR lies in 2^{n-2} ($\sum$ over subsets of #pairs = $C(n,2) \cdot 2^{n-2}$), and so on — "sum of f over all subsets" usually collapses to one closed form per element/pair.
- "Count pairs (i, j) with condition" $\rightarrow$ sort + two pointers, or sweep left $\rightarrow$ right holding earlier elements in a Fenwick over values (inversions pattern).
- Answer is min/max cost with weird operations $\rightarrow$ hunt for an **invariant** first (sum mod k, parity, coloring) — it gives impossibility proofs AND constructions.
- Operations look irreversible $\rightarrow$ **reverse the process** (deletions become insertions, run queries backwards, last move first).
- Optimize over sequences with an exchangeable order $\rightarrow$ **sort by an exchange argument**: comparator "a before b iff ab-combined beats ba-combined" (classic: concatenate numbers into the largest string).

Arrays & queries

- Range ADD, values read only at the end $\rightarrow$ **difference array**: $d[1] += x$, $d[r+1] -= x$, prefix at the end. 2D: $\pm x$ at the 4 corners, 2D prefix.
- "Max number of overlapping intervals" $\rightarrow$ **events**: $+1$ at l , -1 at $r+1$, running maximum of the prefix sum.
- "Sum of max (or max-min) over ALL subarrays" $\rightarrow$ per-element: $a[i]$ is the max of $(i-L) \cdot (R-i)$ subarrays, L/R = nearest greater neighbors via monotonic stack, $O(n)$. **Duplicates**: strict compare on one side, $\geq$ on the other — else equal maxima get counted twice/zero times. Do max and min in separate passes, subtract.
- **# distinct values in range** without Mo: offline, sort queries by r ; sweep r with a Fenwick: $+1$ at position of value, -1 at its previous occurrence; answer = sum over $[l, r]$.
- Online too hard? Almost every query problem collapses **offline**: sort queries by r / by value / by block and sweep. ("edges with weight $\leq X$ connect u, v ?" $\rightarrow$ sort both, grow a DSU.)
- Frequencies + $\sqrt{v}$ threshold: values occurring $> \sqrt{v}$ times number $\leq \sqrt{v}$ ("heavy") — brute-force the heavy ones, bound the light ones per group.
- Divisor-style double loop for $d \{ \text{for } (j = d; j \leq n; j += d) \}$ is $O(n \log n)$, not $O(n^2)$ — harmonic series. Basis of many counting DPs.
- k-th smallest / rank queries on values up to $1e9 \rightarrow$ coordinate-compress, Fenwick over ranks, or ordered-multiset.

Graphs & trees

- **Subtree queries $\rightarrow$ Euler tour**: $\text{subtree}(u)$ is the contiguous range $[\text{in}[u], \text{in}[u] + \text{sz}[u] - 1]$ — Fenwick/segtree on the tour order (the HLD ordering gives this for free).
- Path ADD on a tree, read at the end $\rightarrow$ **tree difference**: $+x$ at u and v , $-x$ at lca , $-x$ at $\text{par}[\text{lca}]$; final value of a node = sum over its subtree (one DFS).
- "Remove edges" / "destroy nodes" queries $\rightarrow$ process **in reverse** with DSU: deletions become additions.
- "Nearest special cell/node" for every position $\rightarrow$ **multi-source BFS** (push all specials at distance 0).
- Edge weights only 0/1 $\rightarrow$ **0-1 BFS** with a deque (push_front on 0-edges) — Dijkstra without the log.
- Collecting sets up a tree (colors in subtree, ...) $\rightarrow$ **small-to-large**: always merge the smaller set into the bigger; total $O(n \log^2 n)$.
- "Answer for EVERY node as root" $\rightarrow$ **rerooting**: one DFS computing down-answers, a second pushing the parent-side answer down (combine prefix/suffix over children).
- Each node has exactly one outgoing edge (functional graph) $\rightarrow$ it's cycles with trees hanging off; k-th successor via binary lifting like `lca.cpp`.

When stuck

- Write the **brute force anyway** ($n \leq 8$): stare at outputs for the pattern, then keep it for stress testing against the real solution with `Other/random.cpp`.
- Fix something and see what's left: the maximum element, the leftmost chosen index, the root — "enumerate the boss, optimize the rest".
- Constraints are a hint (rough map):
 $n \leq 11 \rightarrow O(n!)$ · $n \leq 20 \rightarrow \text{bitmask } 2^n$ · $n \leq 40 \rightarrow \text{meet in the middle}$ ·
 $n \leq 500 \rightarrow O(n^3)$ · $n \leq 5000 \rightarrow O(n^2)$ · $n \leq 2 \cdot 10^5 \rightarrow O(n \log n)$ ·
 $n \leq 10^6 \rightarrow O(n)$ · $n \sim 10^9 \rightarrow O(\sqrt{n})$ / math · $n \sim 10^{18} \rightarrow O(\log n)$: binary search / matrix expo / digit DP.

Number Theory

Divisors

NUMBER OF DIVISORS $d(n)$ + SUM OF DIVISORS $\sigma(n)$ — trial division, $O(\sqrt{n})$

RECALL — both come straight from the prime factorization $n = p_1^{e_1} \cdot p_2^{e_2} \cdot \dots$

$$d(n) = (e_1+1) \cdot (e_2+1) \cdot \dots \quad (\text{pick each prime's exponent independently})$$

$$\sigma(n) = \prod \text{ over } p \text{ of } (1 + p + p^2 + \dots + p^e) \quad (\text{geometric sum per prime})$$

FACTS WORTH REMEMBERING:

- $d(n)$ is ODD $\Leftrightarrow$ n is a PERFECT SQUARE (divisors pair up except $\sqrt{n}$).
- $d(n)$ is small: max 1344 for $n \leq 1e9$, max 103680 for $n \leq 1e18$
 -> "loop over all divisors" is usually fine after factorizing.
- many queries with $n < N$? factorize with spf from linear-sieve.cpp in $O(\log n)$ and apply the same formulas.
- $\sigma(n)$ can be $\sim 6x$ n — near 11's limit for $n \sim 1e18$, careful.
- perfect number: $\sigma(n) = 2n$.

```
long long numberOfDivisors(long long num) {
    long long total = 1;
    for (int i = 2; (long long)i * i <= num; i++) {
        if (num % i == 0) {
            int e = 0;
            do {
                e++;
                num /= i;
            } while (num % i == 0);
            total *= e + 1;           // exponent e contributes (e+1) choices
        }
    }
    if (num > 1) total *= 2;        // leftover prime > sqrt: exponent 1
    return total;
}

long long SumOfDivisors(long long num) {
    long long total = 1;
    for (int i = 2; (long long)i * i <= num; i++) {
        if (num % i == 0) {
            int e = 0;
            do {
                e++;
                num /= i;
            } while (num % i == 0);

            long long sum = 0, pow = 1;    // 1 + p + ... + p^e
            do {
                sum += pow;
                pow *= i;
            } while (e-- > 0);
            total *= sum;
        }
    }
    if (num > 1) total *= (1 + num);    // leftover prime p: (1 + p)
    return total;
}
```

Euler's Phi (Totient)

EULER'S TOTIENT $\phi(n)$ — # of $1..n$ coprime with n

Formula: $\phi(n) = n \cdot \prod_{\text{prime } p|n} (1 - 1/p)$

-> implemented as `res -= res/p` per distinct prime (order doesn't matter).

$\phi(\text{single } n)$: $O(\sqrt{n})$. ϕ for ALL $1..N$: sieve version, $O(N \log \log N)$.

THE FACTS YOU ACTUALLY USE:

- EULER'S THEOREM: $a^{\phi(m)} \equiv 1 \pmod{m}$ when $\gcd(a, m) = 1$.
(generalizes Fermat: $\phi(\text{prime}) = \text{prime} - 1$)
- HUGE EXPONENT REDUCTION (power towers, exponent given as giant number):
 $a^b \equiv a^{(b \bmod \phi(m) + \phi(m))} \pmod{m}$ valid for $b \geq \log_2(m)$,
and this one works EVEN IF $\gcd(a, m) \neq 1$ (the $+\phi(m)$ is what makes it safe – don't drop it unless you know $\gcd = 1$!).
- ϕ is multiplicative: $\phi(a \cdot b) = \phi(a) \cdot \phi(b)$ when $\gcd(a, b) = 1$.
- sum of $\phi(d)$ over all divisors d of $n = n$.
- # of fractions k/n in lowest terms, cycle structures, "count coprime pairs".

```

ll phi(ll n) {                                // O(sqrt n), single value
    ll res = n;
    for (ll p=2; p*p<=n; p++) if (n % p == 0) {
        while (n % p == 0) n /= p;
        res -= res / p;                        // res *= (1 - 1/p)
    }
    if (n > 1) res -= res / n;                  // leftover prime factor > sqrt
    return res;
}

int phiAll[N];                                // phi for every 1..N-1 (sieve style)
void phiSieve() {
    iota(phiAll, phiAll+N, 0);
    for (int i=2; i<N; i++)
        if (phiAll[i] == i)                    // untouched -> i is prime
            for (int j=i; j<N; j+=i)
                phiAll[j] -= phiAll[j] / i;
}

```

Extended GCD & CRT

EXTENDED GCD + MODULAR INVERSE (any mod) + CRT

extgcd(a, b, x, y): returns $g = \gcd(a, b)$ and fills x, y with $a*x + b*y = g$.

WHAT IT UNLOCKS:

- 1) INVERSE MOD ANY m (m need NOT be prime, unlike Fermat in combinatorics.cpp):
 $\text{inv}(a, m)$ valid iff $\gcd(a, m) == 1$.
- 2) LINEAR DIOPHANTINE $a*x + b*y = c$:
 solvable iff $g \mid c$; scale: $x_0 = x*(c/g), y_0 = y*(c/g)$;
 ALL solutions: $x = x_0 + t*(b/g), y = y_0 - t*(a/g), t$ any integer
 (use that to shift x into a wanted range).
- 3) LINEAR CONGRUENCE $a*x \equiv b \pmod{m}$:
 solvable iff $g = \gcd(a, m)$ divides b ; then $x \equiv (b/g)*\text{inv}(a/g, m/g) \pmod{m/g}$,
 i.e. g distinct solutions mod m .
- 4) CRT below.

CRT — combine $x \equiv r_1 \pmod{m_1}$ and $x \equiv r_2 \pmod{m_2}$ into one congruence.

GENERAL version: m_1, m_2 do NOT need to be coprime.

Returns $\{x, \text{lcm}(m_1, m_2)\}$ with $0 \leq x < \text{lcm}$, or $\{-1, -1\}$ if incompatible
 (possible only in the non-coprime case, when $r_1 \not\equiv r_2 \pmod{\gcd}$).

Fold a whole system by chaining: $\text{cur} = \text{crt}(\text{cur}, \text{next}), \dots$

!! pass $0 \leq r < m$ (normalize with $((r \% m) + m) \% m$ first).

Overflow-safe up to $\text{lcm} \sim 9e18$ (`__int128` used at the one risky multiply).

```
ll extgcd(ll a, ll b, ll& x, ll& y) {
    if (!b) { x = 1; y = 0; return a; }
    ll x1, y1, g = extgcd(b, a % b, x1, y1);
    x = y1;
    y = x1 - (a / b) * y1;
    return g;
}

ll inv(ll a, ll m) {
    // inverse of a mod m — REQUIRES gcd(a, m) == 1
    ll x, y;
    extgcd(a, m, x, y);
    return ((x % m) + m) % m;
}

pair<ll, ll> crt(ll r1, ll m1, ll r2, ll m2) {
    ll x, y, g = extgcd(m1, m2, x, y); // m1*x + m2*y = g
    if ((r2 - r1) % g) return {-1, -1}; // no solution
    ll l = m1 / g * m2; // lcm
    ll d = m2 / g;
    ll k = (r2 - r1) / g % d * (x % d) % d; // how many m1-steps to add to r1
    if (k < 0) k += d;
    ll ans = r1 + (ll)((__int128)k * m1 % l); // k*m1 can pass 9e18 mid-multiply
    return {ans % l, l};
}
```

Linear Sieve

LINEAR SIEVE — $O(N)$, gives SMALLEST PRIME FACTOR of every number

$\text{spf}[i]$ = smallest prime factor of i ($\text{spf}[i] == i \iff i$ is prime)

primes = ALL primes $< N$ (complete list, unlike plain sieve variant).

WHY LINEAR: each composite x is crossed out EXACTLY ONCE, by (its smallest prime p) $\cdot (x/p)$ — that's what the ' $\text{primes}[j] \leq \text{spf}[i]$ ' cutoff guarantees.

THE KILLER FEATURE — factorize any $x < N$ in $O(\log x)$:

```
while (x > 1) {
    int p = spf[x], cnt = 0;
    while (x % p == 0) x /= p, cnt++;
    // (p, cnt) is one prime-power of x
}
-> lets you factorize MANY numbers fast (vs trial division  $O(\sqrt{x})$  each).
```

ALSO GOOD FOR multiplicative functions (ϕ , mobius, #divisors) computed for ALL $i < N$ inside this same loop — ask for the extended version if needed.

```
int spf[N];
vector<int> primes;

void LinearSieve() {
    for (int i=2; i<N; i++) {
        if (!spf[i]) {                // nothing crossed i yet -> prime
            spf[i] = i;
            primes.pb(i);
        }

        // cross i * p for primes p <= spf[i]; p becomes the smallest factor of i*p
        for (int j=0; j<primes.size() && i * primes[j]<N && primes[j] <= spf[i]; j++) {
            spf[i*primes[j]] = primes[j];
        }
    }
}
```


Miller-Rabin Primality

MILLER-RABIN — deterministic primality test for ALL 64-bit numbers

isPrime(n) for n up to $\sim 1.8e19$ (unsigned 64-bit), $O(12 * \log n)$ mulmods.

Use when n is way too big for a sieve (e.g. "is $1e18 + 9$ prime?").

RECALL — the test with base a:

write $n-1 = d * 2^r$ (d odd). n passes base a if
 $a^d \equiv 1 \pmod{n}$ or $a^{d*2^i} \equiv n-1$ for some $0 \leq i < r$.
 Every prime passes every base; composites fail at least 3/4 of bases.
 With the FIXED base set {2,3,5,7,...,37} (first 12 primes) the test is
 PROVEN exact for all $n < 3.3e24$ -> covers the whole 64-bit range.
 (The old shorter set {2..13} is only safe below $3.47e12$ — a real WA trap!)

NOTES:

- mulmod uses `__uint128_t` -> GCC/Clang only (fine on Codeforces/ECPC judges).
- `millerRabin(n, a)` with $n \% a == 0$ answers directly (handles small $n / n == a$).
- Everything is unsigned long long (ull) — keep it that way, $n-1$ underflows otherwise... it doesn't since $n \geq 2$, but signed overflow in mulmod WOULD break, hence ull.

RELATED (not included): Pollard's rho for FACTORING 64-bit numbers — ask if needed.

```
ull mulmod(ull a, ull b, ull m) {
    return (__uint128_t)a * b % m;      // 128-bit intermediate, no overflow
}
ull powmod(ull a, ull b, ull m) {
    ull res = 1; a %= m;
    for (; b > 0; b >>= 1) {
        if (b & 1) res = mulmod(res, a, m);
        a = mulmod(a, a, m);
    }
    return res;
}
bool millerRabin(ull n, ull a) {        // does n pass witness a?
    if (n % a == 0) return n == a;      // shares a factor with a -> prime only if n == a
    ull d = n - 1; int r = 0;
    while (d % 2 == 0) d /= 2, r++;      // n-1 = d * 2^r
    ull x = powmod(a, d, n);
    if (x == 1 || x == n - 1) return true;
    for (int i = 0; i < r - 1; i++) {    // keep squaring, look for n-1
        x = mulmod(x, x, n);
        if (x == n - 1) return true;
    }
    return false;                       // composite for sure
}
bool isPrime(ull n) {
    if (n < 2) return false;
    if (n == 2 || n == 3 || n == 5 || n == 7) return true;
    if (n % 2 == 0 || n % 3 == 0) return false;
    // first 12 primes -> deterministic for ALL n < 2^64
    for (ull a : {2ULL, 3ULL, 5ULL, 7ULL, 11ULL, 13ULL, 17ULL, 19ULL, 23ULL, 29ULL, 31ULL, 37ULL})
        if (!millerRabin(n, a)) return false;
    return true;
}
```

Sieve of Eratosthenes

SIEVE OF ERATOSTHENES — $O(N \log \log N)$

prime[i] = is i prime; primes = list of all primes $< N$.

Call Sieve() once at startup.

NOTES:

- Inner loop starts at $i*i$: smaller multiples were already crossed by smaller primes.
i is ll precisely so $i*i$ can't overflow int when N is large ($N > \sim 46341$).
- WARNING: the outer loop stops at $i*i < N$, so `primes` only receives primes up to $\sqrt{N}$! prime[] itself IS correct for all $i < N$. If you need the full primes list, collect it afterwards:
for (int i=2; i<N; i++) if (prime[i]) primes.pb(i);
(or use linear-sieve.cpp, which collects all of them by construction)
- Rough counts: #primes below $1e6 \sim 78498$, below $1e7 \sim 620k$, below $1e8 \sim 5.7M$.

VARIANT — smallest prime factor / $O(\log)$ factorization: see linear-sieve.cpp.

VARIANT — count divisors / sum of divisors for all i: loop $j += i$ from i (not $i*i$)

and do $\text{cnt}[j]++$ — that harmonic double loop is $O(N \log N)$.

```
bool prime[N];
vector<int> primes;

void Sieve()
{
    fill(prime, prime+N, true);
    prime[0] = prime[1] = false;

    for (ll i=2; i*i<N; i++) {          // ll so i*i can't overflow
        if (prime[i]) {
            primes.pb(i);               // WARNING: only primes up to sqrt(N)! see header
            for (ll j=i*i; j<N; j+=i) {
                prime[j] = false;
            }
        }
    }
}
```

Combinatorics

Combinatorics Toolkit

COMBINATORICS (factorials + modular inverse)

nCr / nPr in $O(1)$ per query after $O(n)$ precompute.

REQUIREMENT: mod must be PRIME (inverse uses Fermat: $a^{(mod-2)} = a^{-1}$).

If mod is NOT prime -> use Pascal's triangle (ncr-pascal.cpp)
or factor out common primes manually.

USAGE:

```
Combinatorics C;
C.Init(maxN, 1e9+7); // maxN = largest n you will ever ask about
C.nCr(n, r); C.nPr(n, r);
Also gives you safe modular ops: Add / Sub / Mul / Div / Pow / Inv.
```

RECALL - common formulas:

- Stars and bars: #solutions of $x_1 + \dots + x_k = n$ ($x_i \geq 0$) = $C(n+k-1, k-1)$
- Catalan(n) = $C(2n, n) / (n+1) = C(2n, n) - C(2n, n+1)$
(balanced brackets, binary trees, non-crossing paths)
- Multiset permutations: $n! / (c_1! * c_2! * \dots)$
- Hockey stick: $\sum_{i=r..n} C(i, r) = C(n+1, r+1)$

```
class Combinatorics
{
public:
    ll mod;
    vector<ll> fact, InvFact;

    ll Pow(ll a, ll b) // binary exponentiation,  $O(\log b)$ 
    {
        ll res = 1;
        a %= mod;
        while(b)
        {
            if(b&1) res = (res * a) % mod;
            a = (a * a) % mod;
            b >>= 1;
        }
        return res;
    }

    ll Inv(ll a) {return Pow(a, mod-2);} // ONLY valid for prime mod, gcd(a,mod)=1
    ll Add(ll a, ll b) {return ((a%mod)+(b%mod))%mod;}
    ll Mul(ll a, ll b) {return ((a%mod)*(b%mod))%mod;}
    ll Sub(ll a, ll b) {return ((a-b)%mod + mod)%mod;} // handles negatives correctly
    ll Div(ll a, ll b) {return Mul(a, Inv(b));}

    void Init(ll n, ll x) // call ONCE with the max n needed;  $O(n + \log \text{mod})$ 
    {
        fact.resize(n+1); InvFact.resize(n+1);
        mod = x;

        fact[0] = 1;
        for(int i=1; i<=n; i++) fact[i] = Mul(fact[i-1], i);

        // one Pow call, then walk down: InvFact[i] = InvFact[i+1] * (i+1)
        InvFact[n] = Inv(fact[n]);
        for(int i=n-1; i>=0; i--) InvFact[i] = Mul(InvFact[i+1], i+1);
    }

    ll nPr(ll n, ll r) // ordered pick:  $n! / (n-r)!$ 
    {
        if(n < 0 || r < 0 || n < r) return 0; // out-of-range = 0, so you can call it blindly
        return Mul(fact[n], InvFact[n-r]);
    }

    ll nCr(ll n, ll r) // unordered pick:  $n! / (r! (n-r)!)$ 
    {
        if(n < 0 || r < 0 || n < r) return 0;
        return Mul(nPr(n, r), InvFact[r]);
    }
};
```

Partitions (Pentagonal Theorem)

PARTITION FUNCTION $p(n)$ — Euler's Pentagonal Number Theorem

$p(n)$ = number of ways to write n as a sum of positive integers, ORDER DOESN'T MATTER.

e.g. $p(4) = 5$: 4, 3+1, 2+2, 2+1+1, 1+1+1+1

RECURRENCE (this is the whole theorem):

$p(n) = \text{sum over } k=1,2,3,\dots \text{ of } (-1)^{k-1} * [p(n - g_1(k)) + p(n - g_2(k))]$
 where $g_1(k) = k(3k-1)/2$ and $g_2(k) = k(3k+1)/2$ (generalized pentagonal numbers:
 1, 2, 5, 7, 12, 15, 22, 26, ...). Stop when $g_1(k) > n$.
 Signs alternate in PAIRS: +, +, -, -, +, +, ... (that's the $k \& 1$ check below)

COMPLEXITY: $O(n * \sqrt{n})$ — only $O(\sqrt{n})$ pentagonal numbers are $\leq n$.

$n = 1e5 \rightarrow \sim 3e7$ ops, fine.

This is a SNIPPET (goes inside main/solve). Needs: n , mod .

Base case $p(0) = 1$ (the empty sum).

```
vector<ll> p(n+1);
p[0] = 1;
for (int i=1; i<=n; i++) {
    for (int k=1; ; k++) {
        ll g1 = 1LL * k * (3 * k - 1) / 2;    // pentagonal number
        ll g2 = 1LL * k * (3 * k + 1) / 2;    // its pair

        if (g1 > i) break;                    // both too big -> done with i

        ll term = p[i - g1];
        if (g2 <= i) term = (term + p[i - g2]) % mod;

        if (k & 1) p[i] = (p[i] + term) % mod;    // odd k -> add
        else p[i] = (p[i] - term + mod) % mod;    // even k -> subtract
    }
}
```

nCr - Pascal's Triangle

nCr VIA PASCAL'S TRIANGLE

$c[n][r] = C(n, r) \% \text{mod}$, built with $c[n][r] = c[n-1][r] + c[n-1][r-1]$.

WHEN to prefer this over the factorial version (combinatorics.cpp):

- mod is NOT prime (no modular inverse needed here!)
- or you want plain (unmodded) values for small n (drop the % mod, use ll, n <= 66 fits ll)

COST: $O(N^2)$ time AND memory -> N up to ~3000-4000 max

(ll $c[N][N]$ with $N=4000$ is already ~128 MB; use int if mod < 2^{31} to halve it).

NOTES:

- Entries with $r > n$ stay 0 (globals are zero-initialized) -> safe to read.
- Call pascal() once at startup.

```
ll c[N][N];

void pascal()
{
    for (int i=0; i<N; i++) {
        c[i][0] = c[i][i] = 1;                // C(n,0) = C(n,n) = 1
        for (int j=1; j<i; j++) {
            c[i][j] = (c[i-1][j] + c[i-1][j-1]) % mod; // take or skip the last element
        }
    }
}
```

DP

Digit DP

DIGIT DP — count numbers in a range with a digit property, $O(\text{len} * \text{STATES} * 10)$

THE FRAME (memorize this, it never changes):

```
answer for [L, R] = f(R) - f(L-1)    <- reduces everything to f(X) = count in [0, X]
go(pos, state, tight, started):
    pos      = which digit of X we're placing (left to right)
    tight    = are we still GLUED to X's prefix? (if yes, digit can't exceed X's digit;
               once we place something smaller, tight becomes false FOREVER -> free digits)
    started  = have we placed a nonzero digit yet? (false = still in leading zeros;
               matters when the property cares about actual digits, e.g. "no repeated
               digit", where leading zeros must NOT count as the digit 0)
    state    = THE problem-specific part (example below: digit sum mod K)
```

MEMOIZE only (!tight && started) states: tight paths are a single chain down X (visited once anyway) and not-started is a single all-zeros chain. So dp[][] needs no tight/started dimensions -> small and simple.

EXAMPLE WIRED IN: how many numbers in [0, X] have digit sum divisible by K.

ADAPT: replace `state`'s meaning + the base-case check + the transition line.

Common states: sum mod K, last digit (adjacent constraints), bitmask of used digits, "contains 13" flag, count of some digit — anything small.

PITFALLS:

- memset dp to -1 INSIDE f(X) (each X is a fresh memo — state meanings collide).
- f(L-1) with L = 0 -> f(-1): must return 0 (guard below).
- the empty/never-started path represents the number 0 — decide if the base case should count it (here: yes, sum 0 divides K) and mind it when L = 0.

```
string s;                // X as a string of digits
int K;                   // example parameter
ll dp[20][200];          // dp[pos][state] for the (!tight && started) universe

ll go(int pos, int state, bool tight, bool started) {
    if (pos == sz(s)) return state == 0;          // CHANGE: does `state` satisfy the property?

    if (!tight && started && dp[pos][state] != -1) return dp[pos][state];

    ll res = 0;
    int lim = tight ? s[pos] - '0' : 9;           // glued -> can't exceed X's digit
    for (int dig=0; dig<=lim; dig++) {
        int nstate = (state + dig) % K;           // CHANGE: the transition
        // if the property ignores leading zeros, guard with: started || dig > 0 ? ... : state
        res += go(pos+1, nstate, tight && dig == lim, started || dig > 0);
    }

    if (!tight && started) dp[pos][state] = res;
    return res;
}

ll f(ll x) {
    // count of valid numbers in [0, x]
    if (x < 0) return 0;
    s = to_string(x);
    memset(dp, -1, sizeof dp);
    return go(0, 0, true, false);
}
// answer for [L, R] = f(R) - f(L-1)
```


Longest Increasing Subsequence (LIS)

LIS — Longest Increasing Subsequence, $O(n \log n)$, with full reconstruction

RECALL — the invariant: $d[\text{len}-1]$ = the SMALLEST possible tail value of an increasing subsequence of length len . d is always sorted, so $a[i]$ can extend the longest length whose tail is $< a[i]$ — found by binary search; $a[i]$ then becomes a better (smaller) tail for that length.
(d is NOT an actual subsequence! reconstruction needs the $\text{par}[]$ bookkeeping.)

>>> THE ONE SWITCH <<<

STRICTLY increasing : `lower_bound` (as written)
NON-DECREASING (ties ok): `upper_bound` — that's the entire change
Longest strictly DECREASING: run on negated values (or reversed array).

SNIPPET (inside solve). Needs n , a . Produces:

```
sz(d)    = LIS length
lis      = one actual LIS (values; use indices via par-chain if needed)
lenAt[i] = length of the best increasing subsequence ENDING at i (often useful alone)
```

```
vector<int> d; // d[len-1] = best (smallest) tail for that len
vector<int> lenAt(n), par(n), idxOfLen; // idxOfLen[len-1] = index holding that tail
for (int i=0; i<n; i++) {
    int j = lower_bound(all(d), a[i]) - d.begin(); // upper_bound -> non-decreasing
    if (j == sz(d)) d.pb(a[i]), idxOfLen.pb(i); // extends the longest -> new length
    else d[j] = a[i], idxOfLen[j] = i; // better (smaller) tail for length j+1
    lenAt[i] = j + 1;
    par[i] = j ? idxOfLen[j-1] : -1; // who comes before me in MY subsequence
}

vector<int> lis; // reconstruct by walking parents
for (int cur = idxOfLen.back(); ~cur; cur = par[cur]) lis.pb(a[cur]);
reverse(all(lis));
```

Matrix Exponentiation

MATRIX EXPONENTIATION — linear recurrences in $O(k^3 \log b)$

WHEN: $f(n)$ depends linearly on the previous k terms and n is HUGE ($1e18$).

$f(n) = c_1 * f(n-1) + c_2 * f(n-2) + \dots + c_k * f(n-k) \rightarrow \text{answer} = T^n * \text{base}$

RECIPE (Fibonacci as the example, $f(0)=0, f(1)=1$):

```
T = {{1,1},           // [f(n+1)]   [1 1] [f(n) ]
     {1,0}};          // [f(n)    ] = [1 0] [f(n-1)]
```

```
auto Tn = binPow(T, n); // T^n
```

$f(n) = Tn[1][0] * f(1) + Tn[1][1] * f(0)$ (or read off $Tn[0][1] = f(n)$ directly for fib)

General k : first row = the coefficients $c_1..c_k$, subdiagonal = 1s (shift register).

ALSO WORKS FOR: # of paths of length EXACTLY n in a graph (T = adjacency matrix!),
probabilities/expected values with fixed transitions, "DP with tiny state, giant n ".

!! PITFALL: binPow takes `a` BY REFERENCE and DESTROYS it (squares it in place).

Pass a copy if you need the matrix again: `auto tmp = T; auto Tn = binPow(tmp, n);`

Feasibility: $k^3 * \log(n) \rightarrow k=100, n=1e18$ is $\sim 6e7$ mults, OK.

```
vector<vector<int>> mult(vector<vector<int>>& a, vector<vector<int>>& b) {
    const int n = a.size(), k = a[0].size(), m = b[0].size(); // (n x k) * (k x m)
    vector<vector<int>> ret(n, vector<int>(m));
    for (int i=0; i<n; i++) {
        for (int j=0; j<m; j++) {
            for (int z=0; z<k; z++) {
                ret[i][j] = (ret[i][j] + 1LL * a[i][z] * b[z][j] % mod) % mod;
            }
        }
    }
    return ret;
}

vector<vector<int>> binPow(vector<vector<int>>& a, ll b) { // !! modifies a (see header)
    const int n = a.size(), m = a[0].size();
    vector<vector<int>> ret(n, vector<int>(m));
    for (int i=0; i<n; i++) ret[i][i] = 1; // identity matrix

    while (b) {
        if (b&1) ret = mult(ret, a);
        a = mult(a, a);
        b >>= 1;
    }

    return ret;
}
```

Rerooting

REROOTING — answer "as if rooted at u" for EVERY u, $O(n)$ total

Two DFS passes: (1) normal tree DP computing $down[u]$ = answer inside u's subtree; (2) push a parent-side value $up[v]$ down the tree, combining what the parent knows MINUS v's own contribution.

As written: $ans[u]$ = SUM of distances from u to every node (classic).

down pass: $sz[u]$, $down[u]$ = sum of dist to nodes in u's subtree

push down: $ans[v] = ans[u] - sz[v] + (n - sz[v])$

(crossing edge u-v: v's subtree gets 1 closer, the rest 1 farther)

RECALL — adapting it:

- Works directly whenever the parent-side value is computable by "whole answer at parent MINUS child's contribution" (sum-like combines).
- Combine NOT invertible (max depth, gcd, ...)? Then for each u build PREFIX and SUFFIX combines over its children's values; child v gets $pref[i-1] + suf[i+1] + \text{parent's up-value}$. Same $O(n)$, just more bookkeeping.
- Values on EDGES: fold the edge weight where the +1s are.

namespace Reroot

```
{
    int n; vector<vector<int>>> adj;
    vector<ll> down, ans; vector<int> sz;

    void dfs1(int u, int p) {
        sz[u] = 1, down[u] = 0;
        for (int v : adj[u]) if (v != p) {
            dfs1(v, u);
            sz[u] += sz[v];
            down[u] += down[v] + sz[v];          // every node in v's subtree is 1 farther
        }
    }
    void dfs2(int u, int p) {
        for (int v : adj[u]) if (v != p) {
            // reroot u -> v: v's subtree (sz[v] nodes) gets closer, rest gets farther
            ans[v] = ans[u] - sz[v] + (n - sz[v]);
            dfs2(v, u);
        }
    }

    void solve() {
        // fill n and adj first (0-indexed)
        sz.assign(n, 0), down.assign(n, 0), ans.assign(n, 0);
        dfs1(0, -1);
        ans[0] = down[0];
        dfs2(0, -1);          // now ans[u] is correct for every u
    }
}
```

SOS DP (Sum over Subsets)

SOS DP (sum over subsets) — $f[m]$ = sum of $a[s]$ over all submasks s of m

$O(B * 2^B)$ for B bits — beats the naive 3^B submask loop when you need ALL masks.

Think of it as a B -dimensional prefix sum: one bit-dimension per pass.

USAGE: load point values into f , run the loop, read answers from f (in-place).

RECALL:

- SUPERSET sums: flip the condition — add $f[m \mid 1 \ll b]$ when bit b is NOT set.
- Inverse (recover point values from submask sums / do exact-cover inclusion-exclusion): same loops, SUBTRACT instead of add.
- "count pairs with $a \& b == 0$ ": b must be a submask of $\sim a \rightarrow$ SOS over complements. " $a \& b == a$ " pairs $\rightarrow$ superset sums.
- Single mask only? enumerate its submasks directly in $O(2^{\text{popcount}})$:
 for (int $s = m$; s ; $s = (s - 1) \& m$) { ... use s ... } // add $s=0$ if needed
 Over ALL m together that's the 3^B total — fine for $B \leq \sim 20$ offline.
- max/min instead of $+$: works the same (it's just a different monoid).

namespace SOS

```
{
    int B; // number of bits, masks are 0 .. 2^B - 1
    vector<ll> f; // in: f[mask] = point value; out: f[mask] = sum over submasks

    void sos() {
        for (int b = 0; b < B; b++)
            for (int m = 0; m < (1 << B); m++)
                if (m >> b & 1) f[m] += f[m ^ (1 << b)];
        // superset version: if (!(m >> b & 1)) f[m] += f[m | (1 << b)];
        // inverse: same loops, f[m] -= ... (undoes one bit-dimension per pass)
    }
}
```

Tree Knapsack

TREE KNAPSACK / "merging subtree DPs" — $O(n^2)$ total (NOT n^3 !)

$dp[u][k]$ = best total value of a CONNECTED subgraph that contains u , stays inside u 's subtree, and has exactly k vertices.
(connected + contains $u \iff$ "u plus some children subtree-pieces")

RECALL — the merge (it's just 0/1 knapsack over children):

```
process child v: dp[u][i + j] <- dp[u][i] + dp[v][j]
- i loops DOWNWARD over what u had BEFORE this child -> each child used once
  (same reason 0/1 knapsack loops capacity downward)
- loops are capped by size[u]-so-far and size[v]; that cap is the whole magic:
  each PAIR of vertices (x in earlier part, y in v) is counted once, at their LCA
-> total work = #pairs =  $O(n^2)$ . NEVER loop to  $n$  — that makes it  $O(n^3)$ .
```

SNIPPET (inside solve). Needs: n , adj , a (node values), INF ($\sim 1e18$ with ll dp).

-INF marks "impossible size". Answer: max over $dp[0][k]$ (root=0) for the k you need;
for "best anywhere in the tree", also consider $dp[u][k]$ at every u — any connected subgraph is captured at its TOPMOST node.

VARIANTS with the same skeleton: "keep k edges/vertices, maximize", counting versions (sum instead of max, mod), costs on edges (add edge cost when $j > 0$).

```
vector<vector<ll>> dp(n, vector<ll>(n+1, -INF));
vector<int> size(n);
function<void(int,int)> dfs = [&](int u, int p) {
    dp[u][1] = a[u]; // just u itself
    size[u] = 1;
    for (int v : adj[u]) if (v != p) {
        dfs(v, u);

        // knapsack-merge v into u; i DOWNWARD, both loops size-capped (see header)
        for (int i=size[u]; i;i--) {
            for (int j=1; j<=size[v]; j++) {
                if (dp[u][i] == -INF || dp[v][j] == -INF) continue;
                dp[u][i+j] = max(dp[u][i+j], dp[u][i] + dp[v][j]);
            }
        }
        size[u] += size[v];
    }
};
dfs(0, -1);
```

Graphs & Trees

Articulation Points

ARTICULATION POINTS (cut vertices) — Tarjan, $O(n + m)$

Cut vertex = node whose REMOVAL (with its edges) disconnects the graph.

RECALL — same in/low machinery as bridges, DIFFERENT condition:

non-root u is a cut vertex $\Leftrightarrow$ SOME child v has $low[v] \geq in[u]$
 (v 's subtree can't climb strictly ABOVE $u \rightarrow$ removing u strands it)
 root is special: cut vertex $\Leftrightarrow$ it has ≥ 2 DFS CHILDREN.
 >>> compare: bridge condition was $low[v] > in[u]$ (strict!) on the EDGE. <<<

NOTES:

- parallel edges DON'T matter here (unlike bridges): removing the VERTEX kills all copies anyway, so skipping the parent by node is fine.
Self-loops: just don't add them to adj.
- a graph can have 0 cut vertices (biconnected) — e.g. a simple cycle.
- one node can be "the" cut vertex for many components; count via how many children satisfy the condition (+1 for the rest of the graph) if a problem asks "into how many pieces does removing u split the graph".

CALL: for (int i=0; i<n; i++) if (!in[i]) dfsCut(i, -1);

```
int n, timer = 1, in[N], low[N];
bool isCut[N];
vector<int> adj[N];

void dfsCut(int u, int p) {
    in[u] = low[u] = timer++;
    int children = 0;
    for (int v : adj[u]) if (v != p) {
        if (in[v])
            low[u] = min(low[u], in[v]);           // back edge
        else {
            dfsCut(v, u); children++;
            low[u] = min(low[u], low[v]);
            if (p != -1 && low[v] >= in[u])         // v's subtree is trapped under u
                isCut[u] = true;
        }
    }
    if (p == -1 && children > 1) isCut[u] = true; // root rule
}
```

Bellman-Ford

BELLMAN-FORD — shortest paths WITH negative edges + negative cycle detection, $O(n*m)$

RECALL: relax every edge, $n-1$ rounds. A shortest path has $\leq n-1$ edges, so after round k all shortest paths using $\leq k$ edges are final. If ROUND n STILL relaxes something $\rightarrow$ a negative cycle is reachable from src.

USAGE:

```
vector<ll> dist;
bool hasNegCycle = bellmanFord(n, src, edges, dist);
dist[u] == INF  $\rightarrow$  unreachable.
```

!! the `dist[u] != INF` guard is mandatory: without it $INF + (\text{negative } w)$ "improves" unreachable nodes and everything is garbage.

NOTES:

- returns true = SOME negative cycle reachable from src. Nodes whose real distance is $-\infty$: collect every v relaxed in the extra pass, then BFS/DFS forward from them – everything visited has $dist = -\infty$.
- negative cycle ANYWHERE in the graph (no source): init ALL $dist$ to 0 (equivalent to a virtual source with 0 -edges to everyone) and run the same loop.
- need the cycle itself: remember $par[v]$ on relaxations; from a node relaxed in the extra pass, follow par n times (guaranteed to land ON the cycle), then loop.

```
struct Edge { int u, v; ll w; };
```

```
bool bellmanFord(int n, int src, vector<Edge>& edges, vector<ll>& dist) {
    dist.assign(n, INF);
    dist[src] = 0;

    for (int it=0; it<n-1; it++) {
        bool any = false;
        for (auto& [u, v, w] : edges)
            if (dist[u] != INF && dist[u] + w < dist[v])
                dist[v] = dist[u] + w, any = true;
        if (!any) break; // already stable  $\rightarrow$  no cycle possible either
    }

    for (auto& [u, v, w] : edges) // round n: anything still moving = neg cycle
        if (dist[u] != INF && dist[u] + w < dist[v])
            return true;
    return false;
}
```


Bridges (Tarjan)

BRIDGES (Tarjan, one DFS, $O(n + m)$)

Bridge = edge whose removal disconnects the graph.

RECALL — how it works:

$in[u]$ = DFS discovery time (also serves as "visited": 0 = unvisited, so timer starts at 1)
 $low[u]$ = smallest discovery time reachable from u 's subtree using tree edges
 + AT MOST ONE back edge.
 Edge ($u \rightarrow$ child v) is a bridge $\Leftrightarrow low[v] > in[u]$
 (v 's subtree has NO back edge climbing above v , so cutting the edge separates it.)

SETUP:

$adj[u]$ holds {neighbor, edge id}. Edge id needed because we skip the PARENT EDGE,
 not the parent node $\rightarrow$ parallel edges ($u-v$ twice) are handled correctly:
 a doubled edge is NOT a bridge, and skipping by node would wrongly report it as one.

CALL: for (int $i=0$; $i<n$; $i++$) if (! $in[i]$) tarjan(i , -1); // graph can be disconnected

READ: $isBridge[e]$ per edge id.

NOTE: recursive — for $n \sim 1e5+$ on small stack limits, convert to iterative
 or use the lambda version (see bridge-tree.cpp).

```

vector<pair<int,int>> adj[N]; // {v, edge id}
int n, m, timer = 1, in[N], low[N], isBridge[N];

void tarjan(int u, int pe) { // pe = id of the edge we came from
    in[u] = low[u] = timer++;
    for (const auto& [v, e] : adj[u]) if (e != pe) { // skip parent EDGE (not node!)
        if (in[v])
            low[u] = min(low[u], in[v]); // back edge -> can climb to v's time
        else {
            tarjan(v, e);
            low[u] = min(low[u], low[v]); // child's reach is our reach
            if (low[v] > in[u]) // child can't climb above u -> bridge
                isBridge[e] = true;
        }
    }
}

```

Bridge Tree

BRIDGE TREE (a.k.a. 2-edge-connected component tree), $O(n + m)$

PIPELINE (3 steps below):

- 1) Mark all bridges (Tarjan – see bridges-using-tarjan.cpp for the theory).
- 2) Flood-fill WITHOUT crossing bridges -> $\text{comp}[u]$ = 2-edge-connected component id.
(inside a component, every pair of nodes has 2 edge-disjoint paths)
- 3) Every bridge connects two different components -> those are the TREE edges.

RESULT: `tree` = the bridge tree on `id` nodes. Any cycle of the original graph is squeezed inside a single tree node, and every original bridge = a tree edge.

WHY YOU WANT IT — turns "edge-connectivity" problems into TREE problems:

- # bridges on the path $u \rightarrow v$ = $\text{tree-distance}(\text{comp}[u], \text{comp}[v])$
- min edges to add to make the whole graph 2-edge-connected = $(\text{leaves} + 1) / 2$
- "does removing edge e disconnect u from v " -> is e a bridge on their tree path

SNIPPET (inside main). Nodes 1-indexed in input, converted to 0-indexed here.
Parallel edges are fine (parent skipped by edge id `pe`, not by node).

```
int n, m; cin >> n >> m;
vector<vector<pair<int,int>>> adj(n); // {v, edge id}
vector<pair<int,int>> edges(m);
for (int i=0; i<m; i++) {
    int u, v; cin >> u >> v;
    adj[--u].emplace_back(--v,i);
    adj[v].emplace_back(u,i);
    edges[i] = {u,v};
}

// ----- step 1: mark bridges -----
int timer = 0;
vector<bool> isBridge(m);
vector<int> in(n,-1), low(n); // in = -1 means unvisited
function<void(int,int)> tarjan = [&](int u, int pe) {
    in[u] = low[u] = timer++;
    for (const auto& [v, e] : adj[u]) if (e != pe) {
        if (~in[v])
            low[u] = min(low[u], in[v]); // back edge
        else {
            tarjan(v, e);
            low[u] = min(low[u], low[v]);
            if (low[v] > in[u]) isBridge[e] = true;
        }
    }
};
tarjan(0,-1); // if the graph may be disconnected: loop over all unvisited nodes

// ----- step 2: components = flood fill that never crosses a bridge -----
vector<int> comp(n,-1); int id = 0;
function<void(int)> setComp = [&](int u) {
    comp[u] = id;
    for (const auto& [v, e] : adj[u]) if (!isBridge[e] && comp[v] == -1) {
        setComp(v);
    }
};
for (int i=0; i<n; i++) {
    if (comp[i] == -1) {
        setComp(i);
        id++;
    }
}

// ----- step 3: bridges become the tree edges -----
vector<vector<int>> tree(id);
for (const auto& [u, v] : edges) {
    int cu = comp[u], cv = comp[v];
    if (cu != cv) { // endpoints in different comps <=> this edge is a bridge
        tree[cu].pb(cv);
        tree[cv].pb(cu);
    }
}
// now run LCA / DFS / DP on `tree`; map original node u -> tree node comp[u]
```

Dijkstra

DIJKSTRA — single-source shortest paths, $O((n + m) \log n)$

!! ONLY for NON-NEGATIVE edge weights. One negative edge -> use bellman-ford.cpp.

USAGE: build adj as {v, w} pairs; dijkstra(src, n); read dist[] / rebuild path via par[].
dist[u] == INF -> unreachable.

RECALL — lazy deletion: we push duplicates into the pq instead of decreasing keys.
The line 'if (d > dist[u]) continue;' throws away stale entries — FORGETTING IT
is the classic bug (turns the algorithm exponential on some tests).

VARIANTS:

- multi-source ("nearest shop"): push ALL sources with dist 0 before the loop.
- weights only 0/1 -> 0-1 BFS with a deque (push_front on 0-edges), $O(n+m)$.
- "k-th cheapest path" / state graphs: node = (vertex, state), same code.
- path RECONSTRUCTION: walk par[] back from target, reverse.

```
vector<pair<int,ll>> adj[N];          // {neighbor, weight}
ll dist[N];
int par[N];

void dijkstra(int src, int n) {
    fill(dist, dist+n, INF);
    priority_queue<pair<ll,int>, vector<pair<ll,int>>, greater<>> pq;    // min-heap {dist, node}
    dist[src] = 0; par[src] = -1;
    pq.push({0, src});

    while (!pq.empty()) {
        auto [d, u] = pq.top(); pq.pop();
        if (d > dist[u]) continue;          // stale entry — DO NOT DELETE THIS LINE
        for (auto& [v, w] : adj[u]) {
            if (d + w < dist[v]) {
                dist[v] = d + w;
                par[v] = u;
                pq.push({dist[v], v});
            }
        }
    }
}
```

DSU (Union-Find)

DSU / UNION-FIND (path compression + union by size)

Maintains disjoint sets under "merge two sets" + "which set is u in".

Both ops amortized $\sim O(1)$ (inverse Ackermann).

USAGE:

```
init(n);           // MUST call first (and again for every test case!)
find(u)            // representative of u's set; find(u)==find(v) <=> same set
merge(u, v)        // returns false if already in the same set
len[find(u)]       // size of u's component (only valid at the ROOT!)
```

RECALL — classic uses:

- Kruskal MST: sort edges by weight, add edge iff merge(u,v) succeeds.
- Count components: components -= merge(u,v) succeeded.
- Offline "connectivity after adding edges" queries.
- Cycle detection in undirected graph: merge fails -> edge closes a cycle.
- Bipartite / parity trick: node u -> two nodes (u, u+n) "u is white / u is black".

namespace DSU

```
{
    int sz; vector<int> par, len;    // par = parent pointer, len = component size (valid at root)

    void init(int n)
    {
        sz = n+3;
        par.assign(sz, 0), len.assign(sz, 1);
        iota(all(par), 0);          // everyone is their own parent
    }

    int find(int u) {
        // path compression: hang u (and the whole chain) directly off the root
        return par[u] == u ? u : par[u] = find(par[u]);
    }

    bool merge(int u, int v) {
        u = find(u), v = find(v);
        if (u == v) return false;    // already connected
        if (len[u] < len[v]) swap(u, v); // union by size: small tree under big tree
        par[v] = u, len[u] += len[v];
        return true;
    }
}
using namespace DSU;
```

DSU with Rollback

DSU WITH ROLLBACK — undo merges in reverse order, $O(\log n)$ per op

!! NO path compression (it can't be undone cheaply) — that's why find is

$O(\log n)$, NOT $\sim O(1)$. Union by size keeps the tree shallow; DON'T "optimize".

!! Rollback is LIFO only: you can undo the LAST merges, not an arbitrary one.

USAGE:

```
init(n);
merge(u, v);           // as usual (failed merges are fine)
int t = snapshot();    // remember a point in time
... more merges ...
rollback(t);           // undo everything after the snapshot
```

RECALL — where this shows up:

- Mo with rollback (mo-with-rollback.cpp): add left-side elements, answer, undo.
- OFFLINE dynamic connectivity ("edge exists during time interval $[l, r]$ "): put each edge on the $O(\log q)$ nodes of a segment tree over TIME covering its interval; DFS the tree merging that node's edges, answer queries at leaves, rollback(t) when leaving the node. Handles deletions with a DSU!
- Bipartiteness / parity under edge insert+undo: same trick, store parity too.

namespace DSU

```
{
    vector<int> par, len;
    vector<pair<int,int>> hist;           // (root that got absorbed, root it went under)

    void init(int n) {
        par.assign(n+3, 0), len.assign(n+3, 1);
        iota(all(par), 0);
        hist.clear();
    }

    int find(int u) {                    // NO compression — walk up every time
        while (par[u] != u) u = par[u];
        return u;
    }

    bool merge(int u, int v) {
        u = find(u), v = find(v);
        if (u == v) return false;       // failed merges push nothing (nothing to undo)
        if (len[u] < len[v]) swap(u, v);
        par[v] = u, len[u] += len[v];
        hist.push_back({v, u});
        return true;
    }

    int snapshot() { return hist.size(); }
    void rollback(int t) {               // undo until only t merges remain
        while ((int)hist.size() > t) {
            auto [v, u] = hist.back(); hist.pop_back();
            par[v] = v, len[u] -= len[v];
        }
    }
}
```

using namespace DSU;

Floyd-Warshall

FLOYD-WARSHALL — ALL-pairs shortest paths, $O(n^3)$, n up to ~400-500

Handles negative EDGES fine; detects negative CYCLES (see below).

RECALL — $d[i][j]$ after iteration k = shortest $i \rightarrow j$ path using only intermediate vertices from $\{0..k\}$. $>>> k$ MUST be the OUTERMOST loop $<<<$ (k inside i/j is the classic silent-wrong-answer bug).

SETUP (snippet, inside solve):

```
d[i][i] = 0; d[u][v] = min edge weight (parallel edges: take min!); else INF.
```

NOTES:

- the INF guards below avoid $INF + w$ overflow/poisoning; alternatively use $INF = 1e18/2$ and clamp, but the guards are cleaner.
- NEGATIVE CYCLE: after the loops, $d[x][x] < 0$ for some $x \Leftrightarrow$ neg cycle. Pair (i, j) has dist $-\infty$ iff it goes through one:
exists x with $d[i][x] \neq INF \ \&\& \ d[x][j] \neq INF \ \&\& \ d[x][x] < 0$.
- path reconstruction: keep $nxt[i][j]$ = first hop of best $i \rightarrow j$; on improve:
 $nxt[i][j] = nxt[i][k]$; walk nxt from i until j .
- also good for transitive closure (bitset version) and min/max bottleneck ("minimize the maximum edge"): replace $+$ with $\max$ and $\min$ stays $\min$.

```
ll d[MAXN][MAXN]; // MAXN ~ 500 -> 2 MB of ll, fine
```

```
void floyd(int n) {
    for (int k=0; k<n; k++)
        for (int i=0; i<n; i++) {
            if (d[i][k] == INF) continue; // nothing goes through k from i
            for (int j=0; j<n; j++) {
                if (d[k][j] == INF) continue;
                d[i][j] = min(d[i][j], d[i][k] + d[k][j]);
            }
        }
}
```

Heavy-Light Decomposition (HLD)

HEAVY-LIGHT DECOMPOSITION (HLD)

Answer PATH queries (sum/min/max on the path $u..v$) with a segment tree, and support point/range updates. Each query touches $O(\log n)$ chains, so total $O(\log^2 n)$ per query with a segtree on top.

RECALL — the idea:

- $big[u]$ = "heavy child" = child with the largest subtree.
- $decompose()$ writes nodes in an order where every heavy chain is CONTIGUOUS, $in[u]$ = position of node u in that order, $head[u]$ = top node of u 's chain.
- Walking $u \rightarrow$ root switches chains only $O(\log n)$ times, because every time you leave a chain top, the subtree size at least doubles.

SETUP (order matters):

- 1) build adj, choose root r (usually 0)
- 2) $dfs0(r, -1);$ // sizes, heavy child, par, lvl
- 3) $timer = 0; head[r] = r; decompose(r, -1);$
- 4) build a segtree of size n ; put value of node u at position $in[u]$.
 $tree[pos]$ tells you WHICH node sits at segtree position pos (for initial build).

QUERY PATH $u..v$:

for (auto $[l, r]$: $getRanges(u, v)$) $ans = merge(ans, seg.query(l, r));$
 NOTE: the ranges come in arbitrary order (we jump from both ends) $\rightarrow$
 fine for COMMUTATIVE ops (sum/min/max/gcd). For non-commutative
 (e.g. hashing along the path) you must collect and order them carefully.

EDGE-WEIGHTED version (costs on edges, not nodes) — 2 changes:

- 1) store each edge's weight on the DEEPER endpoint (child), root holds identity
- 2) in $getRanges$, the shared-chain range must EXCLUDE the LCA itself:
 use the commented ``in[u]+1`` line instead (LCA's slot holds the edge ABOVE lca, which is not on the path).

BONUS TRICK (free with this ordering): the subtree of u is the contiguous range

$[in[u], in[u] + sz[u] - 1]$ $\rightarrow$ subtree queries/updates with the same segtree!

Multi-test: reset $timer = 0$ (and clear adj).

```
int n, timer, sz[N], in[N], tree[N], big[N], head[N], par[N], lvl[N];
vector<int> adj[N];

// sizes, heavy child, parent, level
void dfs0(int u, int p) {
    sz[u] = 1;
    big[u] = -1;
    for (int v : adj[u]) if (v != p) {
        par[v] = u;
        lvl[v] = lvl[u] + 1;
        dfs0(v, u);
        sz[u] += sz[v];
        if (big[u] == -1 || sz[v] > sz[big[u]]) big[u] = v; // heaviest child
    }
}

// write nodes to positions: heavy chain first => chain is contiguous in `in` order
void decompose(int u, int p) {
    tree[timer] = u; // tree[pos] = node at segtree position pos
    in[u] = timer++;
    if (~big[u]) {
        head[big[u]] = head[u]; // heavy child continues MY chain
        decompose(big[u], u);
    }

    for (int v : adj[u]) if (v != p && v != big[u]) {
        head[v] = v; // light child starts its own chain
        decompose(v, u);
    }
}

// O(log n) segtree ranges covering the path u..v (inclusive)
vector<pair<int,int>> getRanges(int u, int v) {
    vector<pair<int,int>> ranges;
    while (true) {
        if (head[u] == head[v]) { // same chain -> final contiguous piece, contains the LCA
            if (lvl[u] > lvl[v]) swap(u, v); // u = higher one = the LCA itself
            ranges.emplace_back(in[u], in[v]);
            // EDGE QUERIES: replace the line above with:
            // if (u != v) ranges.emplace_back(in[u] + 1, in[v]); // skip the LCA slot
            break;
        }
        if (lvl[u] > lvl[v]) u = par[u];
        else v = par[v];
    }
    return ranges;
}
```

```

    }

    // jump up from the chain whose head is DEEPER (else we could skip past the LCA)
    if (lvl[head[u]] < lvl[head[v]]) swap(u, v);
    ranges.emplace_back(in[head[u]], in[u]); // whole chain piece: head..u
    u = par[head[u]]; // hop above the chain
}
return ranges;
}

```


LCA (Binary Lifting)

LCA + K-TH ANCESTOR (binary lifting)

$anc[k][v]$ = the 2^k -th ancestor of v . Build $O(n \log n)$, query $O(\log n)$.

RECALL — how it works:

$anc[0][v] = \text{parent}$. $anc[k][v] = anc[k-1][anc[k-1][v]]$ ("jump $2^{(k-1)}$ twice").

We can fill v 's whole column during the DFS because all of v 's ancestors are already fully computed when we reach v .

kth ancestor: decompose k in binary, jump by each set bit.

LCA(u, v): lift the deeper one to the same level; if equal $\rightarrow$ done;

else jump BOTH up by the largest powers that DON'T make them meet;

they end up just below the LCA $\rightarrow$ answer is $anc[0][u]$.

SETUP (sizes here are placeholders — declare AFTER reading n in real code):

$LOG = \lceil \lg(n) \rceil + 2$; // $2^{LOG} > n$ ($n=2e5 \rightarrow LOG=19$)

$anc.assign(LOG, vector<int>(n));$

$lvl[root] = 0$; $anc[0][root] = root$; // root's "parent" = itself, so jumps CLAMP at root

$dfs0(root, -1)$;

HANDY: $dist(u, v) = lvl[u] + lvl[v] - 2 * lvl[lca(u, v)]$

kth node on path $u \rightarrow v$: if $k \leq lvl[u] - lvl[l]$ lift from u , else lift $dist - k$ from v .

NOTE: because root points to itself, kth_anc with $k > \text{depth}$ just returns the root (never crashes) — often exactly what you want.

```
int n, LOG;
vector<vector<int>>>anc(LOG, vector<int>(n)), adj(n);
vector<int> lvl(n);

void dfs0(int u, int p) {
    for (auto v : adj[u]) if (v != p) {
        lvl[v] = lvl[u] + 1;
        anc[0][v] = u;

        for (int k=1; k<LOG; k++)
            anc[k][v] = anc[k-1][anc[k-1][v]];    // 2^k = two jumps of 2^(k-1)

        dfs0(v, u);
    }
}

// jump k edges up from u (clamps at root)
int kth_anc(int k, int u) {
    for (int i=0; i<LOG; i++) if (k >> i & 1) u = anc[i][u];
    return u;
}

int get_lca(int u, int v){
    if (lvl[u] > lvl[v]) swap(u, v);           // make v the deeper one
    v = kth_anc(lvl[v]-lvl[u], v);             // equalize levels
    if (u == v) return u;                      // u was an ancestor of v

    // biggest-first: jump both up while they stay DIFFERENT
    for (int i=LOG-1; i>=0; i--)
        if (anc[i][u] != anc[i][v])
            u = anc[i][u], v = anc[i][v];

    return anc[0][u];                          // now both sit right below the LCA
}
```

MST (Kruskal)

MST — KRUSKAL, $O(m \log m)$

Sort edges ascending, greedily take any edge whose endpoints are in different components (DSU). Self-contained mini-DSU included below.

USAGE:

```
Kruskal::init(n); Kruskal::edges.push_back({w, u, v}); ...
ll total = Kruskal::mst(); // -1 if graph is disconnected
Kruskal::treeEdges // the n-1 chosen {w, u, v} (build adj from it)
```

RECALL — properties that solve problems (often no MST code needed at all):

- CUT property: the lightest edge crossing ANY cut is in some MST.
- MST minimizes the MAXIMUM edge over the path for EVERY pair u, v
 -> "minimize the max edge $u \rightarrow v$ " = sort edges + DSU, stop when u, v connect
 (no tree needed). All-pairs version: path-max on the MST via lca.cpp.
- MAXIMUM spanning tree: sort descending.
- Which edges can be in some MST? Process equal-weight GROUPS together:
 an edge is useless iff its endpoints were already connected BEFORE its group.
- Second-best MST = min over non-tree edges (u, v, w) of $\text{MST} - \text{maxEdgeOnPath}(u, v) + w$.

```
namespace Kruskal
{
    int n; vector<int> par, len;
    vector<array<ll,3>> edges, treeEdges; // {w, u, v}

    void init(int n_) {
        n = n_;
        par.assign(n+3, 0), len.assign(n+3, 1);
        iota(all(par), 0);
        edges.clear(), treeEdges.clear();
    }
    int find(int u) { return par[u] == u ? u : par[u] = find(par[u]); }
    bool merge(int u, int v) {
        u = find(u), v = find(v);
        if (u == v) return false;
        if (len[u] < len[v]) swap(u, v);
        par[v] = u, len[u] += len[v];
        return true;
    }

    ll mst() {
        sort(all(edges)); // by weight (first field)
        ll total = 0;
        for (auto& [w, u, v] : edges)
            if (merge(u, v)) total += w, treeEdges.push_back({w, u, v});
        // n-1 edges <=> connected; anything less means a forest
        return (int)treeEdges.size() == n - 1 ? total : -1;
    }
}
```

Strongly Connected Components (SCC)

SCC + CONDENSATION (Tarjan, one DFS, $O(n + m)$)

SCC = maximal set of nodes in a DIRECTED graph where everyone reaches everyone.

Condensation = shrink each SCC to one node -> always a DAG.

WHAT THIS FUNCTION DOES (in/out parameters — it REPLACES n and adj!):

```
in : n, adj = original directed graph (0-indexed)
out: n, adj = condensation DAG (multi-edges deduped via `seen`)
    compId[u] = component of original node u (size n_original, pre-sized by caller!)
    sccs[c]   = list of original nodes in component c
Keep copies of the originals if you still need them.
```

CALL:

```
vector<int> compId(n); vector<vector<int>>> sccs;
condense(n, adj, compId, sccs);
```

KEY PROPERTY (Tarjan gives it for free):

```
components are produced in REVERSE topological order ->
for every condensed edge cu -> cv: compId[u] > compId[v].
So looping c = 0..n-1 processes SINKS FIRST — perfect for DAG DP where
dp[c] depends on successors (e.g. "best value reachable from c").
Loop c = n-1..0 for sources-first.
```

RECALL — how Tarjan works:

```
low[u] = lowest discovery time reachable from u using tree edges + back edges
to nodes STILL ON THE STACK (inStack check — cross edges to finished SCCs don't count).
When low[u] == in[u], u is the "root" of an SCC -> pop the stack down to u.
```

CLASSIC USES: 2-SAT, "min nodes to start from to reach all" (= #source comps),

"add min edges to make strongly connected" (= max(#sources, #sinks), 1 comp -> 0).

```
void condense(int& n, vector<vector<int>>& adj, vector<int>& compId, vector<vector<int>>& sccs) {
    vector<bool> inStack(n);
    stack<int> st;
    vector<int> in(n), low(n);    // in = discovery time, 0 = unvisited (timer starts at 1)
    int id = 0, timer = 1;

    function<void(int)> tarjan = [&](int u) {
        in[u] = low[u] = timer++;
        st.push(u); inStack[u] = true;
        for (int v : adj[u]) {
            if (!in[v]) {                // tree edge
                tarjan(v);
                low[u] = min(low[u], low[v]);
            } else if (inStack[v]) {      // back/cross edge into current stack
                low[u] = min(low[u], in[v]);
            }
            // else: v's SCC already finished -> ignore
        }

        if (in[u] == low[u]) {            // u is the root of an SCC -> pop it
            vector<int> comp;
            while (!st.empty()) {
                int x = st.top(); st.pop();
                inStack[x] = false;
                comp.emplace_back(x);
                compId[x] = id;
                if (x == u) break;
            }
            sccs.emplace_back(comp);
            id++;
        }
    };

    for (int i=0; i<n; i++) if (!in[i]) tarjan(i);

    // build condensation, dedup parallel edges between the same pair of comps
    set<pair<int,int>> seen;
    vector<vector<int>> nAdj(id);
    for (int u=0; u<n; u++) {
        for (int v : adj[u]) {
            if (compId[u] != compId[v] && seen.find({compId[u], compId[v]}) == seen.end()) {
                nAdj[compId[u]].emplace_back(compId[v]);
                seen.insert({compId[u], compId[v]});
            }
        }
    }

    n = id;    // n and adj now describe the DAG!
    adj = nAdj;
}
```

Topological Sort

TOPOLOGICAL SORT (Kahn / BFS on in-degrees) — $O(n + m)$

Order of a DAG where every edge $u \rightarrow v$ has u BEFORE v .

USAGE:

```
auto order = toposort(n, adj);
if (sz(order) < n) -> THE GRAPH HAS A CYCLE (that's the standard cycle check!)
```

VARIANTS:

- LEXICOGRAPHICALLY SMALLEST order: replace queue with `priority_queue<int, vector<int>, greater<>>` — $O(n \log n + m)$.
- DP over a DAG (longest path, counting paths): process nodes in this order, push dp along edges. Longest path in DAG = classic "critical path".
- DFS alternative: postorder DFS, reverse it (same as SCC condensation order — note `scc.cpp` already hands you a reverse-topological component order for free).

```
vector<int> toposort(int n, vector<vector<int>>& adj) {
    vector<int> indeg(n);
    for (int u=0; u<n; u++)
        for (int v : adj[u]) indeg[v]++;

    queue<int> q; // priority_queue<...> for lexicographic
    for (int i=0; i<n; i++) if (!indeg[i]) q.push(i);

    vector<int> order;
    while (!q.empty()) {
        int u = q.front(); q.pop();
        order.pb(u);
        for (int v : adj[u])
            if (--indeg[v] == 0) q.push(v);
    }
    return order; // size < n <=> cycle exists
}
```

Range Queries

Fenwick Tree (Point Update, Range Query)

FENWICK / BIT — point update, range SUM query. $O(\log n)$ each, tiny constant.

!! 1-INDEXED !! Valid positions are 1..n. Index 0 breaks the loops (i&-i).

USAGE:

```
Fenwick f(n);
f.update(i, x);    // a[i] += x    (to SET a[i]=v: update(i, v - current))
f.query(l, r);     // sum of a[l..r]
```

RECALL — why it works: t[i] stores the sum of the block of length (i & -i)

ending at i. Prefix sum p(i) walks i -> i - (i&-i) (strip lowest bit);

update walks the other way. Range sum = p(r) - p(l-1).

WHEN Fenwick vs segtree: Fenwick only needs an INVERTIBLE op (sum, xor, count).

min/max range queries -> use segtree/sparse table instead.

XOR variant: replace both `+=` with `^=` and the `-` in query with `^` — done.

CLASSIC TRICKS:

- Count inversions: sweep the array, query how many seen values are > a[i].
- "k-th smallest present value": Fenwick over values + binary lifting descent.
- Values up to 1e9? coordinate-compress first (Other/coordinate-compression.cpp).

Range UPDATE + point query -> see the sibling file (difference-array Fenwick).

Range update + range query -> needs 2 Fenwicks or a lazy segtree.

```
class Fenwick {
    int n;
    vector<ll> t;

    ll p(int i) {                                // prefix sum a[1..i]
        ll ret = 0;
        for (; i>0; i -= i&-i) ret += t[i];
        return ret;
    }

public:
    Fenwick(int _n):n(_n) { t.assign(n+1,0LL); }

    ll query(int l, int r) { return p(r) - p(l-1); }

    void update(int i, ll x) { // add x to a[i]
        if (i < 1) return;     // guard: i=0 has i&-i == 0 -> would loop forever
        for (; i<=n; i += i&-i) t[i] += x;
    }
};
```

Fenwick Tree (Range Update, Point Query)

FENWICK / BIT — RANGE update, POINT query (difference-array trick). $O(\log n)$.

!! 1-INDEXED !! Valid positions are 1..n.

USAGE:

```
Fenwick f(n);
f.update(l, r, x); // a[i] += x for ALL i in [l, r]
f.query(i);       // current value of a[i]
```

RECALL — the trick: keep Fenwick over the DIFFERENCE array d ($d[i] = a[i] - a[i-1]$).

"add x on $[l..r]$ " = $d[l] += x$, $d[r+1] -= x$ (two point updates)

"value of $a[i]$ " = prefix sum $d[1..i]$ (one prefix query)

The array starts as all zeros. If you have initial values, either:

- do `update(i, i, a0[i])` for each i , or
- just add `a0[i]` yourself when reading `query(i)`.

```
class Fenwick {
    int n;
    vector<ll> t;
    void u(int i, ll x) { // point-update the underlying diff-array Fenwick
        for (; i<=n; i += i&-i) t[i] += x;
    }
public:
    Fenwick(int _n):n(_n) { t.assign(n+1,0LL); }
    ll query(int i) { // returns a[i] = prefix sum of diffs
        ll ret = 0;
        for (; i>0; i -= i&-i) ret += t[i];
        return ret;
    }
    void update(int l, int r, ll x) { // add x to a[l..r]
        if (l > r || r < 1) return;
        if (l < 1) l = 1;           // clamp instead of crashing
        u(l, x);
        if (r+1 <= n) u(r+1, -x);  // r == n -> no cancel needed
    }
};
```

Lazy Segment Tree

LAZY SEGMENT TREE — RANGE update, range query. $O(\log n)$ each.

0-INDEXED, query/update(l, r) INCLUSIVE. As written: RANGE ADD + RANGE SUM.

USAGE:

```
SegmentTree st(a);          // builds from vector a in O(n)
st.update(l, r, x);         // a[i] += x for all i in [l, r]
st.query(l, r);             // sum a[l..r]
```

RECALL — lazy propagation:

A node's val is always CORRECT for its whole segment; pending updates that its CHILDREN haven't seen yet sit in `lazy`. Before stepping into children (query or update), propagate() pushes the lazy one level down.
change(x, len) = "apply update x to a whole segment of length len".

===== THE 3 SPOTS THAT DEFINE THE TREE (change these per problem) =====

- 1) merge() — how two children combine (here: sum)
- 2) Node neutral — identity of merge (sum->0, min->INF, max->-INF)
- 3) change() — how an update hits a segment + how lazies COMPOSE

===== COMMON VARIANTS =====

(A) range ADD + range SUM (as written):

change: val += x * len; lazy += x;

(B) range ADD + range MIN/MAX:

merge -> min/max, neutral -> INF/-INF,

change: val += x; lazy += x; // NO len — min shifts by x, not x*len

(C) range ASSIGN (a[i] = x) + range SUM:

change: val = x * len; lazy = x; // lazies OVERWRITE instead of add
and DEFAULT_LAZY must be a value that can never be assigned (isLazy already guards this — that's exactly why isLazy exists: 0 may be a legit assignment!)

(D) assign + min/max: change: val = x; lazy = x;

PITFALLS:

- val/lazy are ll here: range-add + sum overflows int almost immediately.
- neutral Node() is returned for out-of-range parts of query — it MUST match merge.
- Multiple lazy types (add AND assign together) need composition rules — derive carefully (assign kills pending add; add on top of assign folds into the assign).

```
struct Node {
    static ll DEFAULT_VALUE, DEFAULT_LAZY;
    ll val, lazy;
    bool isLazy;

    Node(ll x = DEFAULT_VALUE) : val(x), lazy(DEFAULT_LAZY), isLazy(false) {}

    // apply update x to this ENTIRE segment (len = number of leaves under it)
    void change(ll x, int len) {
        val += x * len; // (A) sum. | (B) min/max: val += x | (C) assign+sum: val = x*len
        lazy += x;      // compose with pending lazy.      | (C)/(D) assign: lazy = x
        isLazy = true;
    }
};

ll Node::DEFAULT_VALUE = 0, Node::DEFAULT_LAZY = 0; // neutral val / neutral lazy

class SegmentTree {
private:
    int TREESIZE;
    vector<Node> tree;

    Node merge(const Node &nl, const Node &nr) { return Node(nl.val + nr.val); } // CHANGE

    void Build(const vector<ll> &a, int idx, int lx, int rx) {
        if (lx == rx) {
            if (lx < (int) a.size()) tree[idx] = Node(a[lx]);
            else tree[idx] = Node(); // padding leaves = neutral
            return;
        }

        int mid = (lx + rx) >> 1;
        Build(a, 2 * idx + 1, lx, mid);
        Build(a, 2 * idx + 2, mid + 1, rx);

        tree[idx] = merge(tree[2 * idx + 1], tree[2 * idx + 2]);
    }

    // push this node's pending lazy down ONE level (children lengths differ!)
    void propagate(int idx, int lx, int rx) {
        if (lx == rx || !tree[idx].isLazy) return;

```



```

    int mid = (lx + rx) >> 1;
    tree[2 * idx + 1].change(tree[idx].lazy, mid - lx + 1);
    tree[2 * idx + 2].change(tree[idx].lazy, rx - mid);

    tree[idx].isLazy = false;
    tree[idx].lazy = Node::DEFAULT_LAZY;
}

Node query(int l, int r, int idx, int lx, int rx) {
    propagate(idx, lx, rx); // ALWAYS push before descending

    if (lx >= l && rx <= r) return tree[idx]; // fully inside
    if (rx < l || lx > r) return Node(); // fully outside -> neutral

    int mid = (lx + rx) >> 1;
    Node nLeft = query(l, r, 2 * idx + 1, lx, mid);
    Node nRight = query(l, r, 2 * idx + 2, mid + 1, rx);

    return merge(nLeft, nRight);
}

void update(int l, int r, ll newVal, int idx, int lx, int rx) {
    propagate(idx, lx, rx);

    if (lx >= l && rx <= r) { // fully inside -> lazy-apply and stop
        tree[idx].change(newVal, rx - lx + 1);
        return;
    }
    if (rx < l || lx > r) return;

    int mid = (lx + rx) >> 1;
    update(l, r, newVal, 2 * idx + 1, lx, mid);
    update(l, r, newVal, 2 * idx + 2, mid + 1, rx);

    tree[idx] = merge(tree[2 * idx + 1], tree[2 * idx + 2]);
}

public:
    SegmentTree(const vector<ll> &a) {
        TREESIZE = 1;
        while (TREESIZE < (int) a.size()) TREESIZE <= 1; // pad to power of two
        tree.resize(2 * TREESIZE);
        Build(a, 0, 0, TREESIZE - 1);
    }

    ll query(int l, int r) { return query(l, r, 0, 0, TREESIZE - 1).val; }

    void update(int l, int r, ll newVal) { update(l, r, newVal, 0, 0, TREESIZE - 1); }
};

```

Mo's Algorithm

MO'S ALGORITHM — OFFLINE range queries in $O((n + q) * \sqrt{n} * \text{cost}(\text{add}/\text{del}))$

Use when: no updates, all queries known upfront, and the answer for $[l, r]$

can be MAINTAINED while moving the endpoints one step at a time.

Classics: # distinct values in range, # pairs, mode-ish counts, XXX-frequency stuff.

RECALL — why it's fast:

Sort queries by (block of l , then r). Inside one block R only sweeps $\sim n$ total (amortized), and L jiggles within a block of size SQ per query.

The $\lceil l/SQ \rceil \text{ ? } r \text{ > ... : } r \text{ < ...}$ is the SNAKE/parity order: alternate R 's sweep direction per block $\rightarrow R$ doesn't reset to the left each block ($\sim 2x$ speedup).

SETUP — needs `const int SQ` defined:

$SQ \sim \sqrt{n}$ (450 for $n=2e5$), or tuned: $SQ = \max(1, n / \max(1, (\text{int})\sqrt{q}))$.

!! INDEXING: this template starts $L=1, R=0$ (empty window) $\rightarrow$ queries are

1-INDEXED inclusive. Keep your array 1-indexed to match.

FILL IN (the whole problem lives in these three lambdas):

`add(i)` — element $a[i]$ enters the window (e.g. if `++cnt[a[i]] == 1` `distinct++`);

`del(i)` — element $a[i]$ leaves the window (e.g. if `--cnt[a[i]] == 0` `distinct--`);

`answer()` — current window's answer

Values up to $1e9 \rightarrow$ coordinate-compress first so `cnt[]` can be an array.

ORDER MATTERS below: expand (add) BEFORE shrink (del) so the window never

goes "inside out" ($L > R+1$) — with e.g. `cnt[]--` on absent elements that corrupts state.

If delete is impossible/hard (max-so-far, DSU...) $\rightarrow$ `mo-with-rollback.cpp`.

Path queries on a tree $\rightarrow$ `mo-on-tree.cpp`.

```
struct Query
{
    int l, r, idx; // idx = original position (answers go back in order)
    bool operator<(const Query& other) const {
        if (l/SQ != other.l/SQ) return l/SQ < other.l/SQ; // primary: block of l
        return ((l/SQ) & 1) ? r > other.r : r < other.r; // snake order on r
    }
};

void MO(vector<Query>& queries, vector<int>& ans)
{
    sort(queries.begin(), queries.end());

    auto add = [&](int i) {
        // a[i] enters the window
    };

    auto del = [&](int i) {
        // a[i] leaves the window
    };

    auto answer = [&]() {
        return 0; // current window answer
    };

    int L = 1, R = 0; // empty window
    for (auto [l, r, idx] : queries) {
        while (l < L) add(--L); // expand left
        while (R < r) add(++R); // expand right
        while (L < l) del(L++); // shrink left
        while (r < R) del(R--); // shrink right

        ans[idx] = answer();
    }
}
```

Mo on Tree

MO'S ON A TREE — offline PATH queries $u..v$, $O((n+q) * \sqrt{n})$

Classic: SPOJ COT2 — # distinct values on the path. Needs: $\text{const SQ} (\sim \sqrt{2n})$, LOG.

RECALL — the flattening trick (Euler tour with BOTH stamps):

Every node u is written TWICE: at $\text{in}[u]$ (entering) and $\text{out}[u]$ (leaving).

$\text{flat}[]$ has length $2n$. For a path query $u..v$ (say $\text{in}[u] \leq \text{in}[v]$):

- u is an ancestor of v ($\text{lca} == u$): $\text{window} = [\text{in}[u], \text{in}[v]]$
- otherwise: $\text{window} = [\text{out}[u], \text{in}[v]]$ + add LCA manually

KEY FACT: inside that window, nodes ON the path appear exactly ONCE;

nodes OFF the path appear TWICE (their whole subtree got entered and left).

So $\text{process}()$ toggles: 1st appearance $\rightarrow$ node enters, 2nd $\rightarrow$ node leaves.

That's the freqNode parity check. The LCA is the only path node missing in case 2, hence the manual add/del around answering.

PITFALLS:

- $\text{a}[i]$ values used as indices into cnt arrays $\rightarrow$ COORDINATE-COMPRESS them if they can exceed n (here input is just shifted to 0-based: values assumed $1..n$; replace with compression when values go to $1e9$).
- add/del receive a POSITION in $\text{flat}[]$; $x = \text{flat}[i]$ is the VALUE at that spot.
- Don't forget to define SQ and LOG ($\text{LOG} = \lceil \lg(n)+2 \rceil$).

```
struct Query {
    int l, r, iq, lca;                // window [l,r] in flat[], iq = original index,
                                     // lca = -1 if ancestor case (LCA already inside window)
    bool operator<(const Query& other) const {
        if (l/SQ != other.l/SQ) return l/SQ < other.l/SQ;
        return ((l/SQ)&1) ? r > other.r : r < other.r;    // snake order (see mo.cpp)
    }
};

void solve()
{
    int n, q; cin >> n >> q;
    vector<int> a(n);
    for (int i=0; i<n; i++) cin >> a[i], --a[i];    // values -> 0-based (COMPRESS if large!)
    vector<vector<int>> adj(n);
    for (int i=0; i<n-1; i++) {
        int u, v; cin >> u >> v;
        adj[--u].pb(--v);
        adj[v].pb(u);
    }

    // ---- Euler tour (in & out stamps) + binary lifting for LCA ----
    int timer = 0;
    vector<int> in(n), out(n), lvl(n), flat(2*n), node(2*n); // node[pos] = which node sits there
    vector<vector<int>> anc(LOG, vector<int>(n));
    function<void(int,int)> dfs0 = [&](int u, int p) {
        in[u] = timer;
        node[timer++] = u;
        for (int v : adj[u]) if (v != p) {
            lvl[v] = lvl[u] + 1;

            anc[0][v] = u;
            for (int i=1; i<LOG; i++)
                anc[i][v] = anc[i-1][anc[i-1][v]];

            dfs0(v, u);
        }

        out[u] = timer;
        node[timer++] = u;
    };
    dfs0(0, -1);

    for (int i=0; i<n; i++) flat[in[i]] = flat[out[i]] = a[i]; // value copied to both stamps

    auto kth_ancestor = [&](int k, int u) {
        for (int i=0; i<LOG; i++) if (k>>i&1) u = anc[i][u];
        return u;
    };

    auto get_lca = [&](int u, int v) {
        if (lvl[u] > lvl[v]) swap(u, v);
        v = kth_ancestor(lvl[v]-lvl[u], v);
        if (u == v) return u;
        for (int i=LOG-1; i>=0; i--)
            if (anc[i][u] != anc[i][v])
                u = anc[i][u], v = anc[i][v];
    };
}
```

```

    return anc[0][u];
};

// ---- build the windows (the case split from the header) ----
vector<Query> queries(q);
for (int i=0; i<q; i++) {
    int u, v; cin >> u >> v, --u, --v;
    if (in[u] > in[v]) swap(u, v);          // ensure in[u] <= in[v]
    int lca = get_lca(u, v);
    if (lca == u) {
        queries[i] = {in[u], in[v], i, -1}; // ancestor case: LCA(=u) is in the window
    }
    else {
        queries[i] = {out[u], in[v], i, lca}; // general case: add LCA by hand later
    }
}
sort(all(queries));

vector<int> answer(q), freqNode(n);          // freqNode = appearances of node in window (0/1/2)

// FILL THESE (same roles as plain Mo): maintain answer over the SET of path nodes
auto add = [&](int i) {
    int x = flat[i];
    // x = the VALUE entering, e.g.: if (++cnt[x] == 1) distinct++;
};

auto del = [&](int i) {
    int x = flat[i];
    // x = the VALUE leaving, e.g.: if (--cnt[x] == 0) distinct--;
};

// parity toggle: 1st time node appears -> enters; 2nd time -> cancels out
auto process = [&](int i) {
    int u = node[i];
    if (++freqNode[u] & 1) add(in[u]);
    else del(in[u]);
};

auto get_answer = [&] {
    return 0; // current answer, e.g. distinct
};

int L = 0, R = -1;                          // 0-indexed empty window (flat is 0-indexed!)
for (const auto& [l, r, iq, lca] : queries) {
    while (l < L) process(--L);
    while (R < r) process(++R);
    while (L < l) process(L++);
    while (r < R) process(R--);
    if (~lca) add(in[lca]);                  // LCA joins just for the answer...
    answer[iq] = get_answer();
    if (~lca) del(in[lca]);                  // ...and leaves again (keeps window clean)
}

for (int x : answer) cout << x << '\n';
}

```

Mo with Rollback

MO WITH ROLLBACKS ("Mo, add-only") — SKELETON, fill the marked spots

Use when ADD is easy but DELETE is impossible/ugly:

running max / "max frequency" / merging DSU components / max subarray-ish info.

Complexity: $O(n \cdot \sqrt{q} + q \cdot \sqrt{n})$ adds — no deletes ever happen.

RECALL — how deletes are avoided:

Group queries by BLOCK OF L; inside a block sort by r ASCENDING.

For one block [lBlock, rBlock]:

- R starts at rBlock+1 and only GROWS through the block's queries
-> right-side adds are permanent for the block, never undone.
- The left part (l .. rBlock) is small ($\leq \sqrt{q}$ elements): add it, answer, then ROLLBACK just those adds (undo via a stack of old values).
- Tiny queries fully inside the block ($r \leq rBlock$) -> brute force over $\leq \sqrt{q}$ elements directly, then undo.

Rollback = restore recorded old values in reverse order. NO delete logic needed — that's the whole point (e.g. you can't "un-max" a running maximum, but you CAN restore its previous snapshot).

TO USE: paste into solve(), declare n, q, SQ ($\sim \sqrt{n}$), the array, and
vector<int> answer(q); then fill every `//` hole below.

```
class MO_With_Rollbacks {
    struct Query {
        int l, r, iq;
        bool operator<(const Query& other) const {
            return r < other.r;          // r ascending WITHIN a block (no snake here!)
        }
    };

    struct State {
        // one undo record: whatever add() overwrites, e.g. {int* where, int oldVal}
        // or {value, oldCnt} — enough to restore EXACTLY one add()'s damage
    };

    MO_With_Rollbacks(vector<int>& a, vector<pair<int,int>>& quer) {}

    vector<int> MO()
    {
        int bCnt = (n+SQ-1)/SQ;          // number of 1-blocks
        vector<vector<Query>>> queries(bCnt); // bucket queries by block of 1
        for (int i=0; i<q; i++) {
            int l, r; cin >> l >> r, --l, --r;
            queries[l/SQ].pb({l,r,i});
        }
        for (auto& qu : queries) sort(all(qu));

        int checkpoint, ans;              // ans = CURRENT maintained answer
        stack<State> changes;              // undo log

        auto reset = [&] {
            // wipe the whole structure (called once per block): cnt[]=0, ans=0, clear changes
        };

        auto rollback = [&] {
            // undo everything after `checkpoint`
            while (sz(changes) > checkpoint) {
                auto& c = changes.top(); changes.pop();
                // restore c's old values (reverse order is automatic via the stack)
            }
        };

        auto add = [&] (int idx) {
            // a[idx] enters: PUSH the old state onto `changes` FIRST, then apply.
            // NOTE: `ans` itself must also be restorable -> record it in State too,
            // or recompute it; easiest is to save {oldAns, ...} in every record.
        };

        for (int ib=0; ib<bCnt; ib++) {
            reset();                      // fresh structure for this block

            int lBlock = ib*SQ, rBlock = min(n-1, lBlock+SQ-1);
            int R = rBlock + 1;           // right pointer, only moves forward

            for (const auto& [l, r, iq] : queries[ib]) {
```

```
if (r <= rBlock) {           // tiny query, lives inside the block
    // brute force: for (int j=1; j<=r; j++) add(j);
    answer[iq] = ans;
    // rollback(); (undo ALL of it – set checkpoint=0 before, or loop-pop)
    continue;
}

while (R <= r) add(R++);      // permanent right-side adds (survive the block)
checkpoint = sz(changes);     // <- everything after this will be undone

for (int j=rBlock; j>=1; j--) add(j);  // temporary left part
answer[iq] = ans;
rollback();                    // undo only the left part
}
}
return answer;
};
```

Segment Tree

SEGMENT TREE (recursive) — point update, range query. $O(\log n)$ each.

0-INDEXED. query(l, r) is INCLUSIVE on both ends.

USAGE:

```
SegmentTree st; st.init(n);
st.update(pos, val);      // a[pos] = val (assignment, not +=; change in update() if needed)
st.query(l, r);
```

CHANGE HERE (the only two spots that define "what the tree computes"):

```
1) merge()      - sum / min / max / gcd / max-subarray-struct / ...
2) Node()       - the NEUTRAL element, must satisfy merge(x, neutral) = x:
                  sum -> 0, min -> INF, max -> -INF, gcd -> 0.
                  (query returns Node() for fully-outside ranges, so a wrong neutral = wrong answers!)
Put anything you need inside Node (e.g. {sum, prefix, suffix, best} for max subarray)
and combine the fields in merge().
```

NOTES:

- Size padded up to a power of two; leaves beyond n hold Node() so they're harmless.
- Memory: allocates $4 \cdot \text{TREESIZE}$ but $2 \cdot \text{TREESIZE}$ is enough (TREESIZE already padded) – wasteful but safe; shrink if memory-tight.
- To BUILD from an initial array in $O(n)$: see lazy-seg-tree.cpp's Build (same idea), or just call update n times ($O(n \log n)$, usually fine).

```
class SegmentTree {
    struct Node {
        int val;
        Node(int x = 0) : val(x) {}      // CHANGE: neutral element (0 for sum, INF for min...)
    };

    int TREESIZE;
    vector<Node> tree;

    Node merge(const Node& nl, const Node& nr) { return Node(nl.val + nr.val); } // CHANGE

    Node query(int l, int r, int idx, int lx, int rx) {
        if (rx < l || lx > r) return Node();      // fully outside -> neutral
        if (lx >= l && rx <= r) return tree[idx]; // fully inside -> take node

        int mid = (lx + rx) >> 1;                // children: 2i+1 = [lx,mid], 2i+2 = (mid,rx]
        return merge(query(l, r, 2*idx+1, lx, mid), query(l, r, 2*idx+2, mid+1, rx));
    }

    void update(int pos, int newVal, int idx, int lx, int rx) {
        if (lx == rx) {                          // reached the leaf
            tree[idx] = Node(newVal);             // CHANGE here for a[pos] += x style updates
            return;
        }

        int mid = (lx + rx) >> 1;
        if (pos <= mid) update(pos, newVal, 2*idx+1, lx, mid);
        else update(pos, newVal, 2*idx+2, mid+1, rx);

        tree[idx] = merge(tree[2*idx+1], tree[2*idx+2]); // recompute on the way up
    }

public:
    SegmentTree(){}

    void init(const int& n) {
        TREESIZE = 1;
        while (TREESIZE < n) TREESIZE <= 1;      // pad to power of two
        tree.assign(4 * TREESIZE, Node());        // 2*TREESIZE suffices – see header
    }

    int query(int l, int r) { return query(l, r, 0, 0, TREESIZE-1).val; }

    void update(int pos, int newVal) { update(pos, newVal, 0, 0, TREESIZE-1); }
};
```

Segment Tree (Iterative)

SEGMENT TREE (iterative / "bottom-up") — point update, range query

Same power as the recursive one but ~2-3x faster constant and works with ANY n (no power-of-two padding; memory exactly 2n).

0-INDEXED, query(l, r) INCLUSIVE (converted to half-open [l, r+1) internally).

LAYOUT RECALL: leaves live at tree[n .. 2n-1] (leaf i at n+i),

internal node i has children 2i and 2i+1, parent of i is i/2.

Query climbs from both ends inward, grabbing a node whenever it's "sticking out" of the range (the l&1 / r&1 checks).

CHANGE HERE:

- 1) merge() — sum / min / max / gcd / ...
- 2) Node() — the NEUTRAL element matching merge (sum->0, min->INF, ...)
- 3) update() — currently ASSIGNS; make it `tree[pos].val += val` for increments.

!! CAVEAT — NON-COMMUTATIVE merges (matrices, "function composition", hashes):

this query accumulates the right side in reverse order. Fine for sum/min/max;

for non-commutative ops keep two accumulators (resL, resR) and merge

resL = merge(resL, tree[l++]) / resR = merge(tree[--r], resR), answer merge(resL, resR).

BUILD from array in O(n): fill tree[n+i] = Node(a[i]),

then for (i = n-1; i >= 1; i--) tree[i] = merge(tree[2i], tree[2i+1]).

```
class SegmentTree {
    struct Node {
        int val;
        Node(int x = 0) : val(x) {}      // CHANGE: neutral element
    };

    int n;
    vector<Node> tree;

    Node merge(const Node& nl, const Node& nr) { return Node(nl.val + nr.val); } // CHANGE

    Node queryy(int l, int r) {          // half-open [l, r)
        Node res = Node();
        for (l+=n, r+=n; l < r; l>>=1, r>>=1) {
            if (l&1) res = merge(res, tree[l++]); // l is a right child -> take it, move right
            if (r&1) res = merge(res, tree[--r]);  // r is a right child -> take the one before
        }
        return res;
    }

public:
    SegmentTree(const int& _n){ init(_n); }

    void init(const int& _n) {
        n = _n;
        tree.assign(n<<1, Node()); // exactly 2n nodes
    }

    void update(int pos, int val) {      // a[pos] = val
        pos += n;                        // jump to the leaf
        tree[pos] = Node(val);
        while (pos > 1) {                 // fix ancestors up to the root
            pos >>= 1;
            tree[pos] = merge(tree[pos << 1], tree[pos << 1 | 1]);
        }
    }

    int query(int l, int r) { return queryy(l, r+1).val; } // inclusive wrapper
};
```


Sliding Window Min/Max

SLIDING WINDOW MIN/MAX — monotonic deque, $O(n)$ total

Min of every window of size k . Deque holds INDICES with $a[.]$ strictly increasing front->back, so the front is always the window minimum.
An element killed from the back can never be an answer again (something newer AND smaller-or-equal exists) — that's the whole proof.

USAGE: `windowMin(a, k) -> res[i] = min of  $a[i .. i+k-1]$  (size  $n-k+1$ ).`
For max: flip the comparison (or negate the array).

RECALL:

- DP speedup: `dp[i] = cost[i] + min(dp[i-k .. i-1])` -> keep the deque over `dp` instead of re-querying a segment tree: $O(n)$ instead of $O(n \log n)$.
- Variable-width windows (two pointers): same deque; pop front while its index < current left bound.
- Monotonic STACK cousin — nearest smaller to the left for every i :
while (!st.empty() && a[st.top()] >= a[i]) st.pop(); // top = answer, then push i
Gives "span until a smaller element" -> largest rectangle in histogram.

```
vector<ll> windowMin(vector<ll>& a, int k)
{
    int n = a.size();
    vector<ll> res;
    deque<int> dq; // indices, a[dq] increasing
    for (int i = 0; i < n; i++) {
        while (!dq.empty() && a[dq.back()] >= a[i]) dq.pop_back(); // > for max
        dq.push_back(i);
        if (dq.front() <= i - k) dq.pop_front(); // fell out of window
        if (i >= k - 1) res.push_back(a[dq.front()]);
    }
    return res;
}
```

Sparse Table

SPARSE TABLE — STATIC array, $O(n \log n)$ build, $O(1)$ range query

No updates allowed! If you need updates -> segment tree.

USAGE: SparseTable st(a); st.query(l, r); // 0-indexed, inclusive

RECALL — how it works:

$t[k][i]$ = merge of the 2^k elements starting at i .

Query $[l..r]$: take the two 2^k blocks covering the range from each end — they OVERLAP in the middle, which is why the op must be IDEMPOTENT:

```
>>> merge must satisfy  $f(x, x) = x$ : min / max / gcd / AND / OR — YES
>>> sum / xor / count — NO! (overlap counts twice)
    for sums use prefix sums; or query the disjoint  $O(\log n)$  way.
```

CHANGE HERE: `merge` (currently min).

$_\lg(x) = \text{floor}(\log_2(x))$ (GCC builtin). Build memory: $O(n \log n)$ ints.

```
class SparseTable {
private:
    vector<vector<int>> t;
    int n, LOG;
    int merge(int e, int o) { return min(e, o); } // CHANGE: min/max/gcd/AND/OR only!

public:
    SparseTable(vector<int>& a) {
        n = a.size(), LOG = 32 - __builtin_clz(n); // = floor(log2(n)) + 1
        t.resize(LOG, vector<int>(n));

        for (int i=0; i<n; i++) t[0][i] = a[i]; // level 0 = the array itself

        for (int k=1; (1<<k)<=n; k++) {
            for (int i=0; i+(1<<k)<=n; i++) {
                // block of  $2^k$  = two adjacent blocks of  $2^{k-1}$ 
                t[k][i] = merge(t[k-1][i], t[k-1][i+(1<<(k-1))]);
            }
        }

        int query(int l, int r) { // 0-indexed inclusive, needs  $l \leq r$ 
            int k = __lg(r-l+1);
            return merge(t[k][l], t[k][r-(1<<k)+1]); // two overlapping blocks
        }
};
```

Sqrt Decomposition

SQRT DECOMPOSITION — $O(\sqrt{n})$ per query/update, dead simple to bend

Split the array into blocks of size $SQ \sim \sqrt{n}$; keep an aggregate per block.

Query $[l, r]$: whole blocks inside use the aggregate, the two ragged ends brute force.

WHEN over a segtree (which is $O(\log n)$ and usually better for plain sum/min):

- the block aggregate is something a segtree can't merge easily:
SORTED block (-> "count elements > x in range" via binary search per block),
frequency map per block, "best pair inside block", etc.
- lazy-style tricks: tag whole blocks, rebuild only ragged blocks (rebuild = $O(SQ)$).
- you're out of time to debug a segtree – this is 15 lines and hard to get wrong.

As written: point ASSIGN + range SUM, 0-indexed inclusive.

CHANGE HERE: what $b[]$ stores + how update/query maintain it (see WHEN above).

```
int SQ;
vector<ll> a, b;                                // b[k] = aggregate (sum) of block k

void build(int n) {                             // call after filling a; O(n)
    SQ = (int)sqrtl(n) + 1;
    b.assign(n / SQ + 1, 0);
    for (int i=0; i<n; i++) b[i / SQ] += a[i];
}

void update(int i, ll x) {                      // a[i] = x
    b[i / SQ] += x - a[i];                     // fix the aggregate, not rebuild
    a[i] = x;
}

ll query(int l, int r) {                       // sum of a[l..r]
    ll sum = 0;
    for (int i=l; i<=r; ) {
        if (i % SQ == 0 && i + SQ - 1 <= r) { // block fully inside -> take aggregate
            sum += b[i / SQ];
            i += SQ;
        }
        else sum += a[i++];                   // ragged edge -> element by element
    }
    return sum;
}
```

Strings

Binary Trie (XOR)

BINARY TRIE — numbers as bit-strings (MSB first). THE tool for XOR problems.

insert / del / maxXor in $O(\text{MAXBITS})$ each.

CHANGE HERE: MAXBITS = highest bit index.

29 covers values $< 2^{30}$. For long long values change: `int` -> `ll`, `1<<i` -> `1LL<<i`,
`MAXBITS` -> 59 (or 62), and ``bool c = n>>i&1`` stays fine.

RECALL — why a trie: walking from the top bit, at each level you can CHOOSE the child that makes the current answer-bit 1 (greedy is optimal because a higher bit beats all lower bits combined).

`freq[c]` = how many stored numbers go through child `c` (supports duplicates; `del()` frees a branch when its `freq` hits 0, so ``child != NULL`` == "branch alive").

CLASSIC USES:

- max XOR pair in an array: insert all, `maxXor(a[i])` over `i`
- max XOR subarray: prefix XORs + this
- online "insert numbers / query best xor with `x`" sets

PITFALL: `del(n)` assumes `n` IS in the trie (like `trie.cpp`). `maxXor` assumes the trie is non-empty.

```
class BinaryTrie {
    static const int MAXBITS = 29;          // CHANGE: 29 -> values < 2^30; 59+11 for big values
    class Node {
    public:
        Node* child[2];
        int freq[2];                          // freq[c] = numbers passing through child c
        Node() {
            memset(child,0,sizeof child);
            memset(freq,0,sizeof freq);
        }
    };
    Node* root = new Node();

public:
    void insert(int n) {
        Node* cur = root;
        for (int i=MAXBITS; i>=0; i--) {
            bool c = n>>i&1;
            if (cur->child[c] == NULL) cur->child[c] = new Node();
            cur->freq[c]++;
            cur = cur->child[c];
        }
    }

    // unwind the path, decrement freqs, free dead branches
    void del(Node* cur, int n, int i) {
        if (i == -1) return;
        bool c = n>>i&1;
        del(cur->child[c], n, i-1);

        cur->freq[c]--;

        if (!cur->freq[c]) {
            delete cur->child[c];
            cur->child[c] = NULL;
        }
    }

    void del(int n) {                          // !! only call if n is currently in the trie
        del(root,n,MAXBITS);
    }

    // max of (x XOR y) over all stored y — greedy: prefer the OPPOSITE bit
    int maxXor(int x) {
        Node* cur = root;
        int res = 0;
        for (int i=MAXBITS; i>=0; i--) {
            bool c = x>>i&1;
            if (cur->child[!c]) {                // opposite bit available -> answer bit = 1
                res |= 1<<i;                    // (1LL<<i for the ll version!)
                cur = cur->child[!c];
            }
            else cur = cur->child[c];            // forced to match -> answer bit = 0
        }
        return res;
    }
};
```

String Hashing

STRING HASHING (polynomial, DOUBLE hash) — substring compare in $O(1)$

$\text{hash}(s[0..i]) = s[0] \cdot b^i + s[1] \cdot b^{i-1} + \dots + s[i] \pmod{m}$, for TWO

independent (base, mod) pairs $\rightarrow$ collision chance $\sim 1/1e18$, safe to treat

"hashes equal" as "strings equal".

USAGE:

```
init(); // ONCE globally (fills power tables, needs N >= max length)
Hash h(s); // O(|s|)
h.get(l, r); // pair hash of s[l..r], 0-INDEXED INCLUSIVE, O(1)
Hash h2(t); h.get(..) == h2.get(..) // cross-string compare works (same bases)
```

RECALL — how get() works: $h1[r]$ holds hash of prefix $0..r$; subtract prefix

$0..l-1$ SHIFTED UP by $b^{(r-l+1)}$ to align, i.e. $\text{hash}(l..r) = h[r] - h[l-1] \cdot b^{\text{len}}$.

NOTES / PITFALLS:

- chars are mapped to (value + 1) so no character is 0 — otherwise "a" and "aa" would collide (leading zeros vanish).
- $b1, b2$ must be $>$ alphabet size. For mixed case/digits use e.g. 131, 137.
- vs anti-hash hacks (Codeforces open hacking): randomize the bases at runtime, e.g. $b1 = \text{rand}(256, \text{mod1}-2)$ using `Other/random.cpp`.
For ECPC onsite nobody hacks you — fixed bases are fine.

CLASSIC USES:

- compare any two substrings / check palindrome (hash the reversed string too)
- longest common prefix of two suffixes via binary search + get()
- count distinct substrings of each length, string matching without KMP

```
const int b1 = 41, b2 = 37, mod1 = 1e9+7, mod2 = 1e9+9;
int pw1[N], pw2[N]; // pw[i] = b^i (precomputed powers)

void init() {
    pw1[0] = pw2[0] = 1;
    for (int i=1; i<N; i++) {
        pw1[i] = 1LL * pw1[i-1] * b1 % mod1;
        pw2[i] = 1LL * pw2[i-1] * b2 % mod2;
    }
}

class Hash
{
    int n;
    vector<int> h1, h2; // prefix hashes

public:
    Hash(const string& s): n(s.size()), h1(n), h2(n) {
        h1[0] = h2[0] = (unsigned char)s[0] + 1; // +1: never map a char to 0!
        for (int i=1; i<n; i++) {
            int d = (unsigned char)s[i] + 1;
            h1[i] = (1LL*h1[i-1]*b1%mod1 + d)%mod1;
            h2[i] = (1LL*h2[i-1]*b2%mod2 + d)%mod2;
        }
    }

    pair<int,int> get(int l, int r) { // hash of s[l..r], 0-indexed inclusive
        int x = (h1[r] - (1LL*(l?h1[l-1]:0)*pw1[r-l+1]%mod1) + mod1) % mod1;
        int y = (h2[r] - (1LL*(l?h2[l-1]:0)*pw2[r-l+1]%mod2) + mod2) % mod2;
        return {x, y};
    }
};
```

KMP (Prefix Function)

PREFIX FUNCTION + KMP MATCHING — $O(n)$

$pi[i]$ = length of the longest PROPER prefix of $s[0..i]$ that is also a suffix
of $s[0..i]$ ("border" length). e.g. "abca**b**" -> $pi = 0\ 0\ 0\ 1\ 2$

RECALL — the loop: to extend position i , try the previous border $j = pi[i-1]$;
if $s[i]$ doesn't match $s[j]$, FALL BACK to the border-of-the-border $pi[j-1]$,
repeat. Borders of borders enumerate ALL borders — that's the whole trick.

PATTERN MATCHING: build pi on $p + \# + t$ ($\#$ = any char not in either
string — it stops borders from crossing the boundary). $pi == |p|$ -> full match.

CLASSIC FACTS (asked constantly):

- smallest PERIOD of $s = n - pi[n-1]$; it tiles s perfectly iff $n \% period == 0$.
- all borders of s : $pi[n-1]$, $pi[pi[n-1]-1]$, ... (chain down to 0).
- count occurrences of every prefix / build prefix automaton — extensions of pi .
- "shortest string with s as prefix and suffix"-type tricks: think borders.

```
vector<int> prefixFunction(const string& s) {
    int n = sz(s);
    vector<int> pi(n);
    for (int i=1; i<n; i++) {
        int j = pi[i-1];
        while (j && s[i] != s[j]) j = pi[j-1];    // best border so far
        if (s[i] == s[j]) j++;                  // fall back through smaller borders
        pi[i] = j;
    }
    return pi;
}

// all START positions (0-indexed) of pattern p inside text t
vector<int> matches(const string& t, const string& p) {
    int m = sz(p);
    vector<int> pi = prefixFunction(p + '#' + t), res;
    for (int i=m+1; i<sz(pi); i++)
        if (pi[i] == m) res.pb(i - 2*m);        // i is in "combined" coords -> shift back
    return res;
}
```

Manacher (Palindromes)

MANACHER — ALL palindromic substrings in $O(n)$

$d1[i]$ = # of ODD palindromes centered at i (= radius incl. center)

longest odd palindrome centered at i has length $2*d1[i] - 1$.

$d2[i]$ = # of EVEN palindromes centered BETWEEN $i-1$ and i

longest even one there has length $2*d2[i]$.

e.g. "abaa": $d1 = 1\ 2\ 1\ 1$, $d2 = 0\ 0\ 0\ 1$ ("aa" centered between index 2,3)

RECALL — same z-box idea as z-function: keep the rightmost known palindrome

$[l, r]$; inside it, position i mirrors to $l+r-i$, so $d[i]$ starts from the

mirror's value (capped by the box), then extends by brute force. $O(n)$ amortized.

INSTANT ANSWERS AFTER $O(n)$ BUILD:

- longest palindromic substring: max over $2*d1[i]-1$ and $2*d2[i]$.

- is $s[l..r]$ a palindrome, $O(1)$: ($len = r-l+1$)

len odd : $d1[(l+r)/2] \geq (len+1)/2$

len even: $d2[(l+r+1)/2] \geq len/2$

- total # of palindromic substrings = $\text{sum}(d1) + \text{sum}(d2)$.

```
vector<int> d1, d2;
```

```
void manacher(const string& s) {
    int n = sz(s);
    d1.assign(n, 0); d2.assign(n, 0);

    for (int i=0, l=0, r=-1; i<n; i++) {
        // odd lengths
        int k = (i > r) ? 1 : min(d1[l+r-i], r-i+1); // mirror value, capped
        while (0 <= i-k && i+k < n && s[i-k] == s[i+k]) k++; // extend
        d1[i] = k--;
        if (i+k > r) l = i-k, r = i+k; // new rightmost palindrome
    }

    for (int i=0, l=0, r=-1; i<n; i++) {
        // even lengths
        int k = (i > r) ? 0 : min(d2[l+r-i+1], r-i+1);
        while (0 <= i-k-1 && i+k < n && s[i-k-1] == s[i+k]) k++;
        d2[i] = k--;
        if (i+k > r) l = i-k-1, r = i+k;
    }
}
```


Multiset Hashing

MULTISET HASHING — order-INDEPENDENT hashing (compare multisets in $O(1)$)

Question it answers: "do these two collections contain exactly the same elements with the same multiplicities?" — regardless of order.
(e.g. is substring A an ANAGRAM of substring B; are two subtrees equal as bags)

IDEA: give every DISTINCT VALUE a random 5-dimensional vector (Hash()).

The hash of a multiset = plain SUM of the vectors of its elements.

add element $v \rightarrow cur = cur + H[v]$

remove element $v \rightarrow cur = cur - H[v]$

Sum is commutative $\rightarrow$ order can't matter. Forging a collision needs to hit

5 independent random $\sim 1e9$ coordinates $\rightarrow$ probability ~ 0 .

USAGE:

```
map<int, Hash> H; auto getH = [&](int v){ if (!H.count(v)) H[v] = Hash(); return H[v]; };
Hash cur(0,0,0,0,0); // empty multiset (explicit zeros – Hash() is RANDOM!)
cur = cur + getH(v); // v enters
cur = cur - getH(v); // v leaves
sliding window / two windows  $\rightarrow$  compare with ==
```

NOTES:

- No mod anywhere: coords stay $< \sim 2e14$ for $2e5$ elements, fits ll fine.
- PITFALL: default constructor Hash() generates a NEW RANDOM vector – that's the per-value generator. The "zero" hash must be built as Hash(0,0,0,0,0).
- Works inside prefix sums too: $pre[i] = pre[i-1] + H[a[i]] \rightarrow$ hash of any subarray = $pre[r] - pre[l-1] \rightarrow$ "are these two ranges permutations of each other" in $O(1)$.

```
mt19937_64 rng(chrono::steady_clock::now().time_since_epoch().count());
template <typename T> T rand(T l, T r) { return uniform_int_distribution<T>(l, r)(rng); }
```

```
struct Hash {
    ll a, b, c, d, e;
    Hash(ll _a, ll _b, ll _c, ll _d, ll _e):a(_a),b(_b),c(_c),d(_d),e(_e){}
    Hash() { // RANDOM vector – use once per distinct value
        a = rand(1, (int)1e9);
        b = rand(1, (int)1e9);
        c = rand(1, (int)1e9);
        d = rand(1, (int)1e9);
        e = rand(1, (int)1e9);
    }
    Hash operator+(const Hash& other) const {
        return Hash(a+other.a,b+other.b,c+other.c,d+other.d,e+other.e);
    }
    Hash operator-(const Hash& other) const {
        return Hash(a-other.a,b-other.b,c-other.c,d-other.d,e-other.e);
    }
    bool operator==(const Hash& other) const {
        return a==other.a&&b==other.b&&c==other.c&&d==other.d&&e==other.e;
    }
};
```

Trie

TRIE (prefix tree) — insert / search / delete words, $O(|\text{word}|)$ each

Node bookkeeping (this is what you usually query):

count = how many inserted words PASS THROUGH this node
 -> "how many words start with prefix p": walk p, return count (0 if walk dies)
 isEnd = how many inserted words END exactly here (int, so duplicates are counted)

CHANGE HERE: N (alphabet size) + baseChar — 'a' for lowercase, 'A', '0' for digits...

PITFALLS:

- del(s) assumes s WAS inserted (call search first if unsure) — deleting an absent word corrupts counts. It removes ONE occurrence and frees dead branches.
- Memory: $\sim(\text{total inserted chars}) \text{ nodes} * (26 \text{ pointers} + 2 \text{ ints}) \approx 220\text{B}/\text{char}$.
 1e6 total chars is fine; 1e7 is not — use arrays-of-int trie instead.
- Multi-test: no clear() here — recreate the Trie object (leaks, but CP doesn't care) or add a recursive delete.

COMMON EXTENSIONS (one-liners to add):

- countPrefix(p): walk p, return cur->count
- countWord(s): walk s, return cur->isEnd
- XOR problems on numbers -> binary-trie.cpp

```
class Trie {
    const static int N = 26;           // CHANGE: alphabet size
    const static int baseChar = 'a';   // CHANGE: first char of alphabet
    class Node {
    public:
        Node* child[N];
        int count, isEnd;
        Node():isEnd(false), count(0) {
            memset(child, 0, sizeof child);
        }
    };
    Node* root = new Node();

public:
    void insert(const string& s) {
        Node* cur = root;
        for (char ch : s) {
            int c = ch - baseChar;
            if (cur->child[c] == NULL) cur->child[c] = new Node();
            cur = cur->child[c];
            cur->count++;           // one more word passes through
        }
        cur->isEnd++;
    }

    // internal: unwind the path, decrement counts, free branches that died
    void del(Node* cur, const string& s, int i) {

        if (i == (int)s.size()) {
            cur->isEnd--;
            return;
        }

        int c = s[i] - baseChar;
        Node* nxt = cur->child[c];
        del(nxt, s, i+1);

        nxt->count--;

        if (!nxt->count) {           // no word passes here anymore -> prune
            delete nxt;
            cur->child[c] = NULL;
        }
    }

    void del(const string& s) {       // !! only call if s is currently in the trie
        del(root, s, 0);
    }

    bool search(const string& s) {    // whole-word lookup (prefix: return cur->count instead)
        Node* cur = root;
        for (char ch : s) {
            int c = ch - baseChar;
            if (cur->child[c] == NULL) return false;
            cur = cur->child[c];
        }
        return cur->isEnd;
    }
};
```

Geometry

Closest Pair of Points

CLOSEST PAIR OF POINTS — $O(n \log n)$ sweep line. Returns SQUARED distance.

Needs P / dist2 from geometry-basics.cpp. $n \geq 2$. Duplicates $\rightarrow$ returns 0.

RECALL — the sweep: process points left to right, keep an "active strip" of points whose x is within $D = \sqrt{\text{best}}$ of the current point, ordered by y (the set). For each new point only candidates with $|dy| \leq D$ matter — and a $D \times 2D$ box can only hold $O(1)$ points that are pairwise $\geq D$ apart, so the inner loop is amortized constant. best only shrinks $\rightarrow$ total $O(n \log n)$.

The only double-ish moment is computing $D = \text{ceil}(\sqrt{\text{best}})$ to bound the strip; all COMPARISONS stay in integers (dist2), so precision can't hurt correctness — D is just a (safe, rounded-up) search radius.

```
11 closestPair(vector<P> p) {
    sort(all(p));                                // by x (then y) – P's operator<
    set<pair<ll, ll>> strip;                       // {y, x} of points in the active strip
    ll best = LLONG_MAX;
    int j = 0;                                    // strip's left edge (index into p)
    for (int i=0; i<sz(p); i++) {
        ll D = (ll)ceil(sqrtl((long double)best));
        while (j < i && p[i].x - p[j].x >= D) // too far left to ever matter  $\rightarrow$  evict
            strip.erase({p[j].y, p[j].x}), j++;

        for (auto it = strip.lower_bound({p[i].y - D, LLONG_MIN});
             it != strip.end() && it->first <= p[i].y + D; ++it)
            best = min(best, dist2({it->second, it->first}, p[i]));

        strip.insert({p[i].y, p[i].x});
    }
    return best;                                // squared! sqrt only for output
}
```

Geometry Basics

GEOMETRY BASICS — integer-only building blocks (no doubles until you must print)

Coordinates up to 1e9 are safe: cross/dot/dist2 peak at ~8e18, fits ll.

THE ONE TOOL THAT SOLVES 90% OF GEOMETRY — the cross product:

```
cross(a, b) = a.x*b.y - a.y*b.x      (of vectors)
cross(o, a, b) = cross(a-o, b-o), its SIGN tells you where b is
    relative to the line o->a:
    > 0 -> b is to the LEFT (counter-clockwise turn o->a->b)
    < 0 -> b is to the RIGHT (clockwise turn)
    = 0 -> o, a, b COLLINEAR
|cross(o,a,b)| = 2 * area of triangle o,a,b.
```

dot(a, b) > 0: angle < 90; = 0: perpendicular; < 0: angle > 90.

PITFALLS:

- NEVER compare double distances – compare SQUARED ll distances (dist2).
- Read ints into ll immediately; subtraction before multiplication is what keeps everything in range.
- Polygon functions expect vertices in order (CW or CCW), no repeated last point.

```
struct P {
    ll x, y;
    P operator-(const P& o) const { return {x - o.x, y - o.y}; }
    P operator+(const P& o) const { return {x + o.x, y + o.y}; }
    bool operator<(const P& o) const { return x < o.x || (x == o.x && y < o.y); }
    bool operator==(const P& o) const { return x == o.x && y == o.y; }
};

ll cross(P a, P b) { return a.x*b.y - a.y*b.x; }
ll cross(P o, P a, P b) { return cross(a-o, b-o); }      // sign: see header
ll dot(P a, P b) { return a.x*b.x + a.y*b.y; }
ll dist2(P a, P b) { ll dx=a.x-b.x, dy=a.y-b.y; return dx*dx + dy*dy; } // SQUARED distance
int sgn(ll v) { return (v > 0) - (v < 0); }

// is p on segment a-b (endpoints included)? collinear + inside the bounding box
bool onSeg(P a, P b, P p) {
    return cross(a, b, p) == 0
        && min(a.x, b.x) <= p.x && p.x <= max(a.x, b.x)
        && min(a.y, b.y) <= p.y && p.y <= max(a.y, b.y);
}

// do segments a-b and c-d share ANY point (crossing, touching, or overlapping)?
bool segIntersect(P a, P b, P c, P d) {
    int d1 = sgn(cross(c, d, a)), d2 = sgn(cross(c, d, b));
    int d3 = sgn(cross(a, b, c)), d4 = sgn(cross(a, b, d));
    if (d1 != d2 && d3 != d4) return true; // proper crossing: each seg separates the other's ends
    return onSeg(a,b,c) || onSeg(a,b,d) || onSeg(c,d,a) || onSeg(c,d,b); // touch / overlap cases
}

// TWICE the signed area of any simple polygon (shoelace).
// result > 0 -> vertices are CCW, < 0 -> CW. Real area = abs(...) / 2.0
// (keep the x2 version to stay in integers – area itself can be x.5)
ll area2(vector<P>& poly) {
    ll s = 0;
    for (int i=0, n=sz(poly); i<n; i++)
        s += cross(poly[i], poly[(i+1)%n]);
    return s;
}

// point vs simple polygon (ANY simple polygon, convex or not), O(n):
// returns 0 = outside, 1 = on the boundary, 2 = strictly inside (ray casting)
int inPolygon(vector<P>& poly, P p) {
    int n = sz(poly), cnt = 0;
    for (int i=0; i<n; i++) {
        P a = poly[i], b = poly[(i+1)%n];
        if (onSeg(a, b, p)) return 1;
        cnt ^= ((p.y < a.y) - (p.y < b.y)) * cross(p, a, b) > 0; // edge crosses the upward ray?
    }
    return cnt ? 2 : 0;
}
```

Other

Bit Tricks

BIT TRICKS — counting set bits in ranges + XOR prefix, O(1) each

(the one-line bit identities live in NOTES.md; this file is the actual code)

countSetBits(l, r, bit): how many numbers in [l, r] have `bit` set.

RECALL — bit b cycles with period 2^{b+1} : first half off, second half on.

Count in [0, x] = full periods * half + the partial period's overlap; subtract prefixes.

Total set bits in a range: loop this over all 60 bits (still $\sim O(60)$).

Typical use: sum/xor contributions bit by bit ("count pairs whose AND has bit b", ...).

xorUpTo(n) = $1 \oplus 2 \oplus \dots \oplus n$. RECALL — it's periodic with period 4:

$n \% 4 == 0 \rightarrow n, \quad == 1 \rightarrow 1, \quad == 2 \rightarrow n+1, \quad == 3 \rightarrow 0$

XOR of any range [l, r] = xorUpTo(r) ^ xorUpTo(l-1) (prefix trick, xor is its own inverse).

```

11 onesUpTo(11 x, int bit) {           // # of v in [0, x] with `bit` set
    if (x < 0) return 0;
    11 len = 1LL << (bit + 1), half = 1LL << bit;
    return (x / len) * half + max(0LL, x % len - half + 1);
}

11 countSetBits(11 l, 11 r, int bit) { // # of v in [l, r] with `bit` set
    return onesUpTo(r, bit) - onesUpTo(l - 1, bit);
}

11 xorUpTo(11 n) {                     // 1 ^ 2 ^ ... ^ n    (n >= 0)
    switch (n % 4) {
        case 0: return n;
        case 1: return 1;
        case 2: return n + 1;
        default: return 0;
    }
}

```

Coordinate Compression

COORDINATE COMPRESSION — map big values to 0..k-1, order preserved. $O(n \log n)$

WHEN: values up to $1e9/1e18$ but only n of them, and you need them as ARRAY

INDICES (Fenwick, segtree, cnt[] in Mo's, ...). Comparisons still work because the mapping is monotonic: $a[i] < a[j]$ before $\Leftrightarrow$ after.

SNIPPET (inside solve). After it:

```
a[i] = compressed rank in [0, k)   where k = sz(d)
d[a[i]] = the ORIGINAL value (keep d around to answer/output real values!)
rank of arbitrary x: lower_bound(all(d), x) - d.begin()
```

NOTE: compressing MULTIPLE arrays together (queries + array values)?

Concatenate everything into d before sort+unique, then rank each against d.

```
vector<int> d = a;
sort(d.begin(), d.end());
d.resize(unique(d.begin(), d.end()) - d.begin()); // sorted distinct values
for (int i = 0; i < n; ++i) {
    a[i] = lower_bound(d.begin(), d.end(), a[i]) - d.begin(); // value -> its rank
}
```


Ordered Multiset (PBDS)

ORDERED SET / MULTISSET (GNU pb_ds) — std::set + ORDER STATISTICS in $O(\log n)$

The two extra superpowers over std::set:

s.order_of_key(x) -> how many elements are STRICTLY LESS than x (x's rank)
 s.find_by_order(i) -> iterator to the i-th smallest (0-indexed); *it to read
 -> "k-th smallest" and "count elements < x" without a Fenwick.

ordered_set<T> : unique elements, everything works like std::set.

ordered_multiset<T> : duplicates via the less_equal comparator HACK — the tree thinks equal elements are "different", which BREAKS some members:

!! find(x) -> ALWAYS returns end() (never use)
 !! erase(x) -> silently erases NOTHING (never use)
 !! lower_bound / upper_bound SWAP meanings (lower_bound acts like upper_bound & vice versa)
 OK: insert, order_of_key (still = count of elements < x), find_by_order.

ERASE ONE COPY of x (the only safe way):

```
auto it = ms.find_by_order(ms.order_of_key(x)); // iterator to first occurrence of x
if (it != ms.end() && *it == x) ms.erase(it); // guard: x might be absent
```

Include the two headers below (they're not in bits/stdc++.h on some setups).

```
#include <ext/pb_ds/assoc_container.hpp>
#include <ext/pb_ds/tree_policy.hpp>
using namespace __gnu_pbds;

template<class T>
using ordered_set = tree<T, null_type, less<T>, rb_tree_tag, tree_order_statistics_node_update>;

template<class T>
using ordered_multiset = tree<T, null_type, less_equal<T>, rb_tree_tag,
tree_order_statistics_node_update>;
```

Random (RNG)

RANDOM — the ONLY rng you should use in CP (never rand()/srand: weak + 32767 cap on MSVC)

rand(l, r) -> uniform integer in [l, r], BOTH INCLUSIVE. Type follows the args:

rand(1, 100) -> int rand(1LL, (1L)1e18) -> ll (don't mix int and ll args!)

OTHER USES:

```
shuffle(all(v), rng);           // kill adversarial input order (anti-quicksort tests)
rng() by itself                  // raw random 64-bit number
```

WHEN randomness saves you:

- random bases/keys for hashing (Strings/ hashing, multiset-hashing)
- randomized algorithms: pick random pivot/sample, "guess a likely answer" tricks
- stress-testing: generate small random tests to diff brute force vs solution

```
mt19937_64 rng(chrono::steady_clock::now().time_since_epoch().count());
template <typename T> T rand(T l, T r) { return uniform_int_distribution<T>(l, r)(rng); }
```